

# Graph Neural Networks

## Introduction

Acknowledgement: Jure Leskovec, Stanford University

# Overview

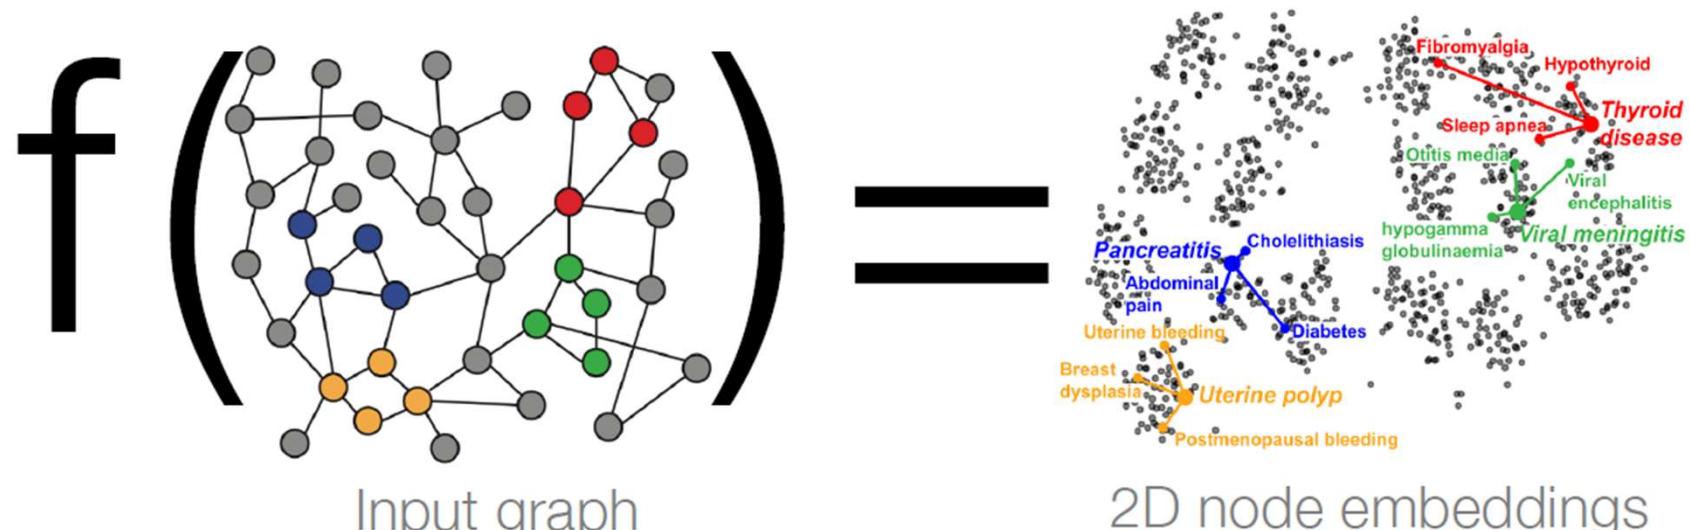
- Shallow Encodings - Limitations
- Deep Learning for Graphs
- Graph Convolution Networks

# Overview

- Shallow Encodings - Limitations
- Deep Learning for Graphs
- Graph Convolution Networks

# Node Embeddings

- **Intuition:** Map nodes to  $d$ -dimensional embeddings such that similar nodes in the graph are embedded close together

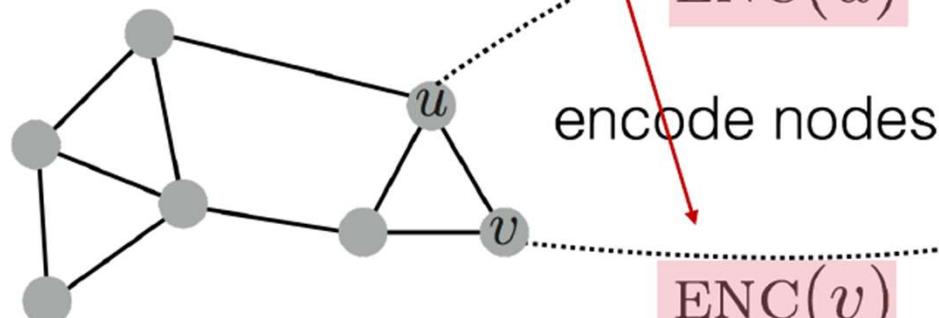


How to learn mapping function  $f$ ?

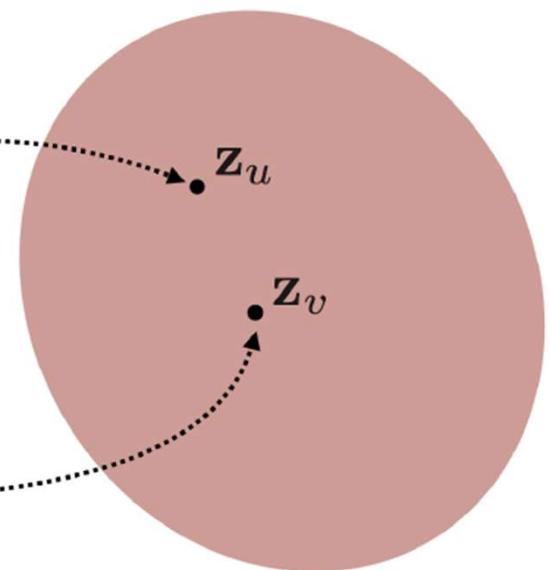
# Node Embeddings

Goal:  $\text{similarity}(u, v) \approx \mathbf{z}_v^T \mathbf{z}_u$

Need to define!



Input network



$d$ -dimensional  
embedding space

# Two Key Components

- **Encoder:** maps each node to a low-dimensional vector

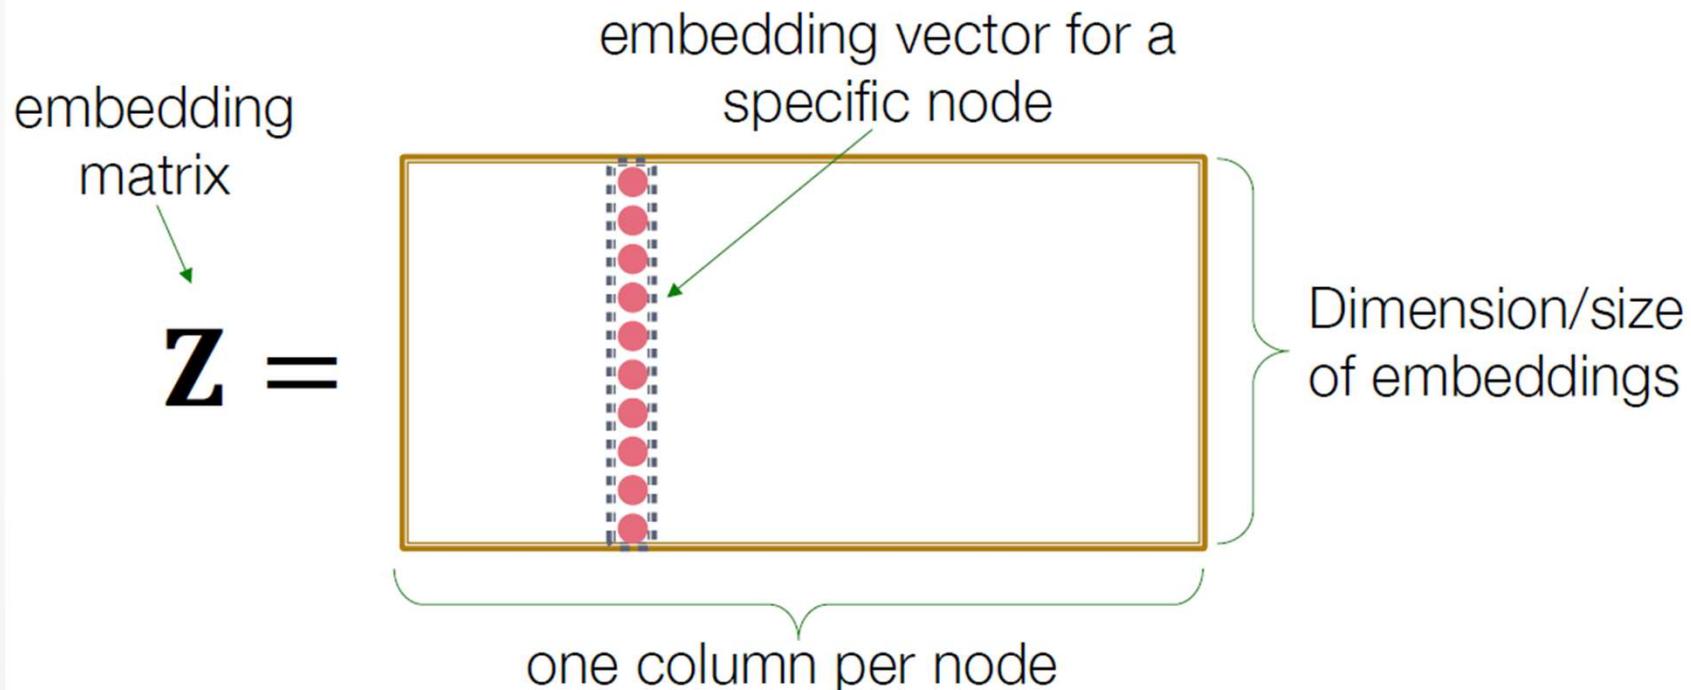
$\text{ENC}(v) = \boxed{\mathbf{z}_v}$  *d*-dimensional embedding  
node in the input graph

- **Similarity function:** specifies how the relationships in vector space map to the relationships in the original network

$\text{similarity}(u, v) \approx \mathbf{z}_v^T \mathbf{z}_u$  **Decoder**  
Similarity of  $u$  and  $v$  in the original network dot product between node embeddings

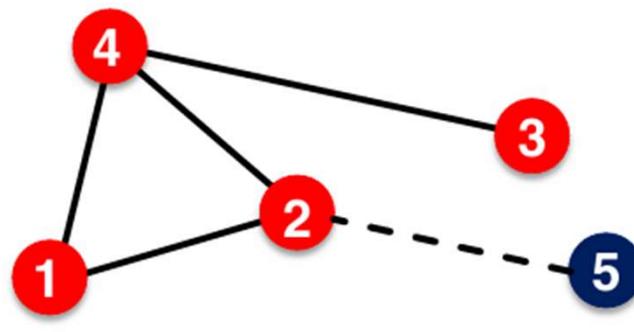
# Shallow Encodings

Simplest encoding approach: **encoder is just an embedding-lookup**



# Shallow Encodings Limitations - 1

- Cannot obtain embeddings for nodes not in the training set



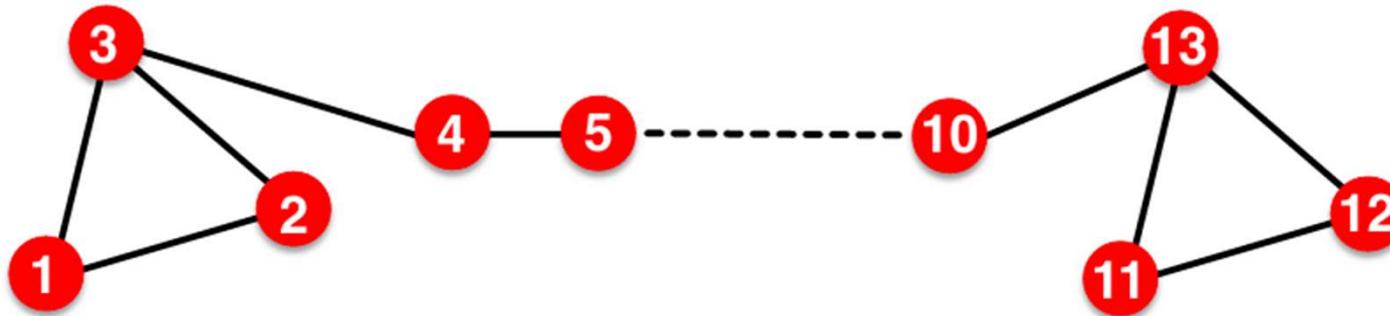
Training set

A newly added node 5 at test time  
(e.g. new user in a social network)

Cannot compute its embedding  
with DeepWalk / node2vec. Need to  
recompute all node embeddings.

# Shallow Encodings Limitations - 2

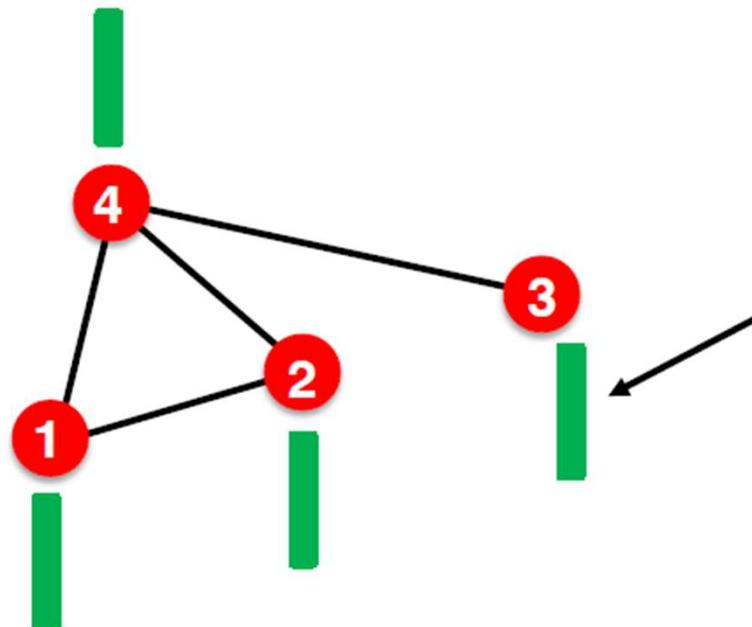
- Cannot capture **structural similarity**:



- Node 1 and 11 are **structurally similar** – part of one triangle, degree 2
- However, they have very **different** embeddings
  - It's unlikely that a random walk will reach node 11 from node 1
- **DeepWalk and node2vec do not capture structural similarity**

# Shallow Encodings Limitations - 3

- Cannot utilize node, edge and graph features



**Feature vector**  
(e.g. protein properties in a  
protein-protein interaction graph)

DeepWalk / node2vec  
embeddings do not incorporate  
such node features

Solution to these limitations: Deep Representation Learning and Graph Neural Networks

# Shallow Encodings Limitations - 4

- Every node has its own unique embedding
- No sharing of parameters between nodes
- $O(|V|d)$  parameters needed

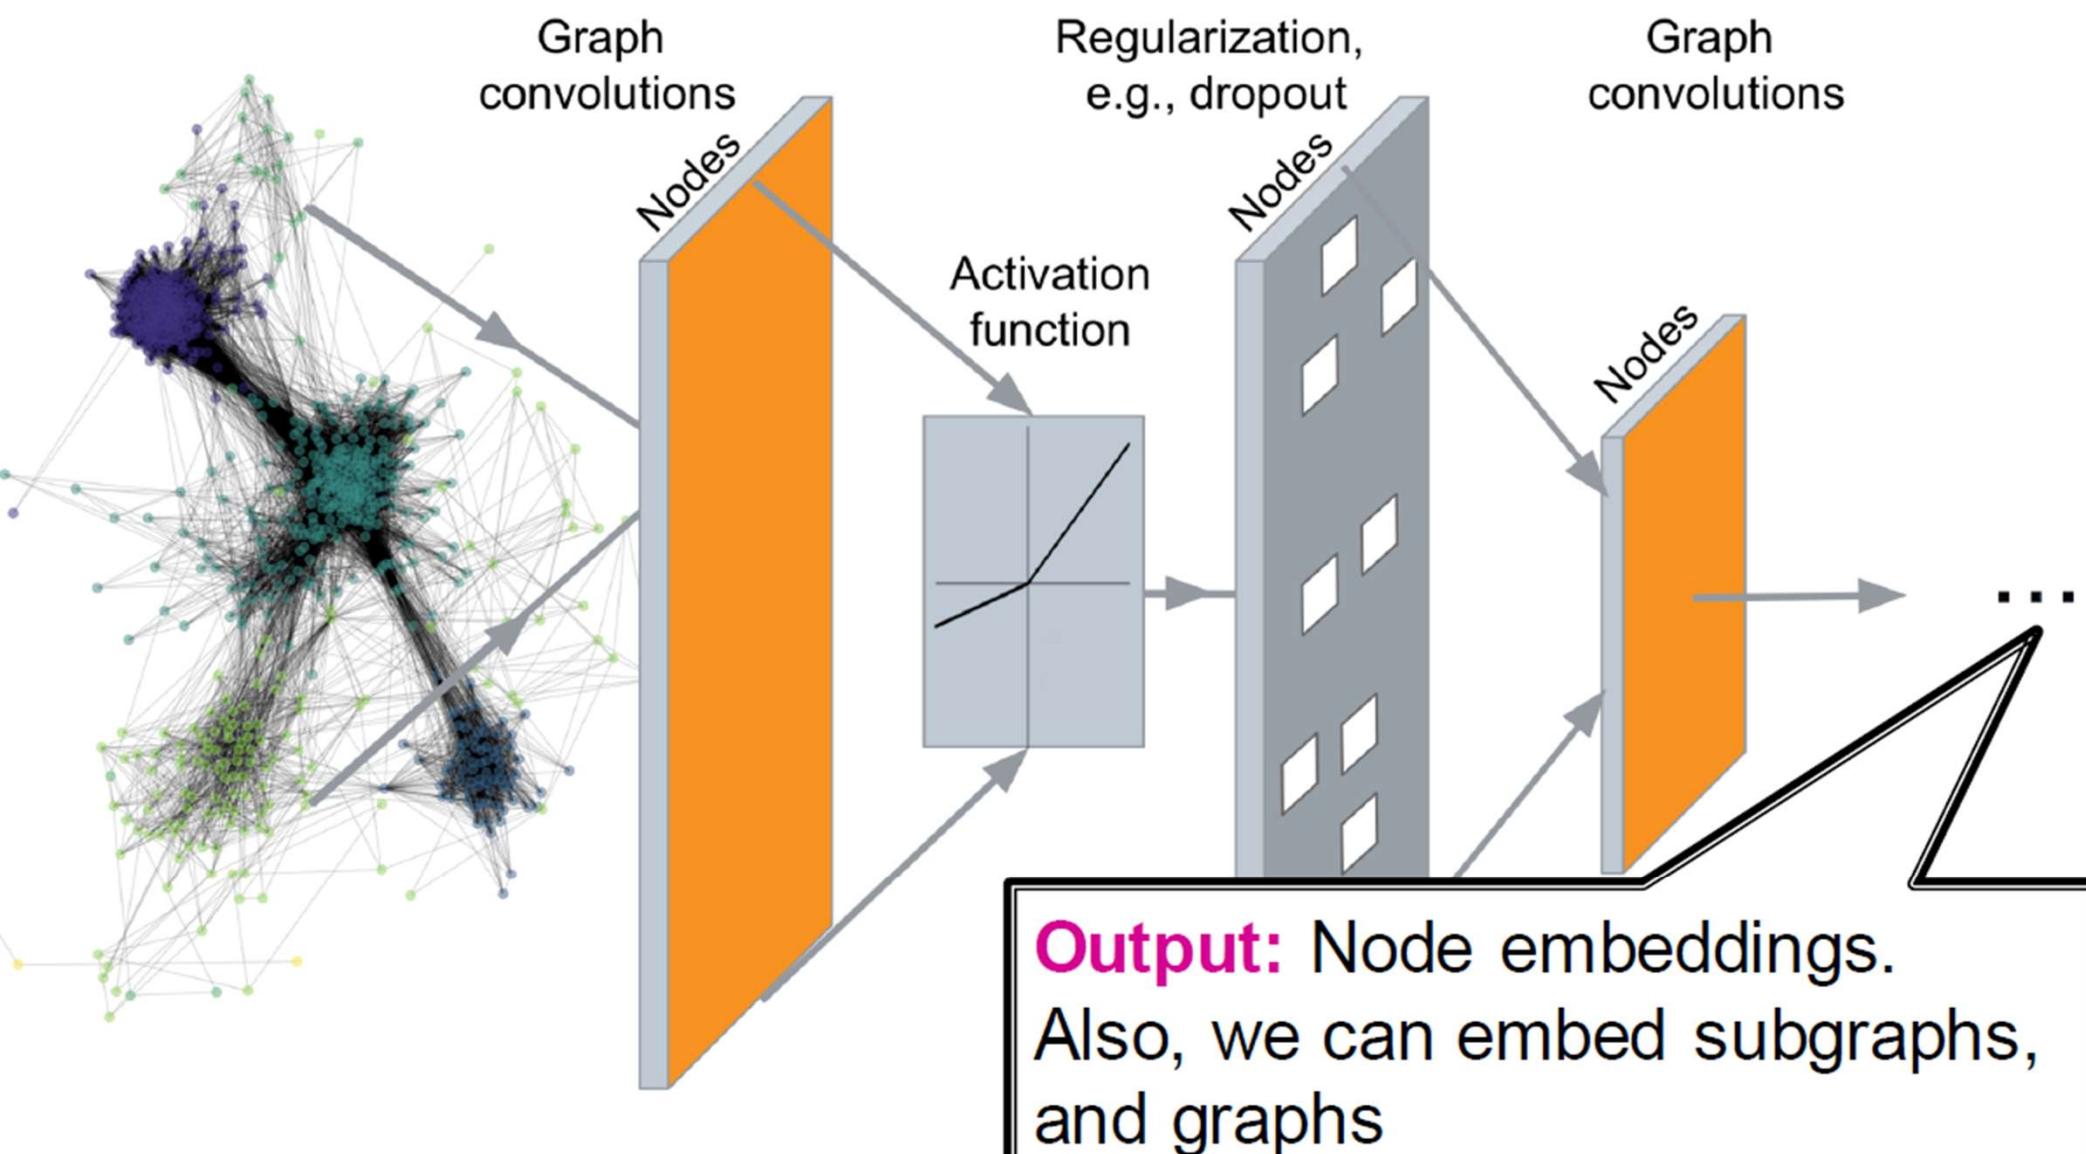
# Deep Encoders

- Deep Encoders based on Graph Neural Networks

$\text{ENC}(v) =$  **multiple layers of  
non-linear transformations  
based on graph structure**

- Deep Encoders can be combined with Graph Similarity functions
-

# Deep Graph Encoders



# Tasks on Networks

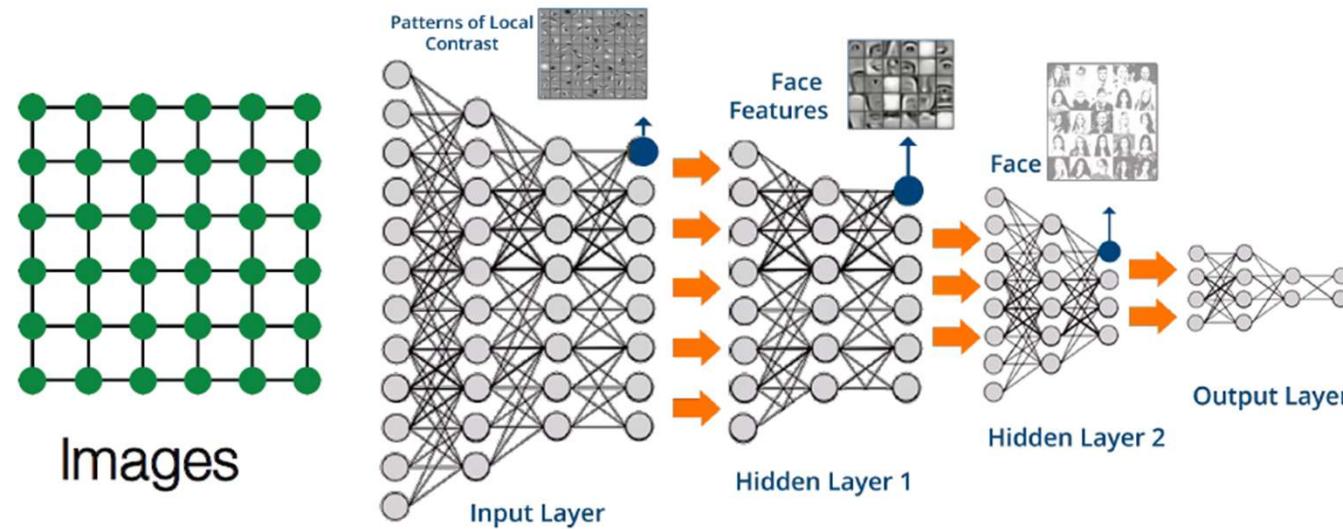
## Tasks we will be able to solve:

- Node classification
  - Predict the type of a given node
- Link prediction
  - Predict whether two nodes are linked
- Community detection
  - Identify densely linked clusters of nodes
- Network similarity
  - How similar are two (sub)networks

# Overview

- Shallow Encodings - Limitations
- Deep Learning for Graphs
- Graph Convolution Networks

# Modern ML Toolbox



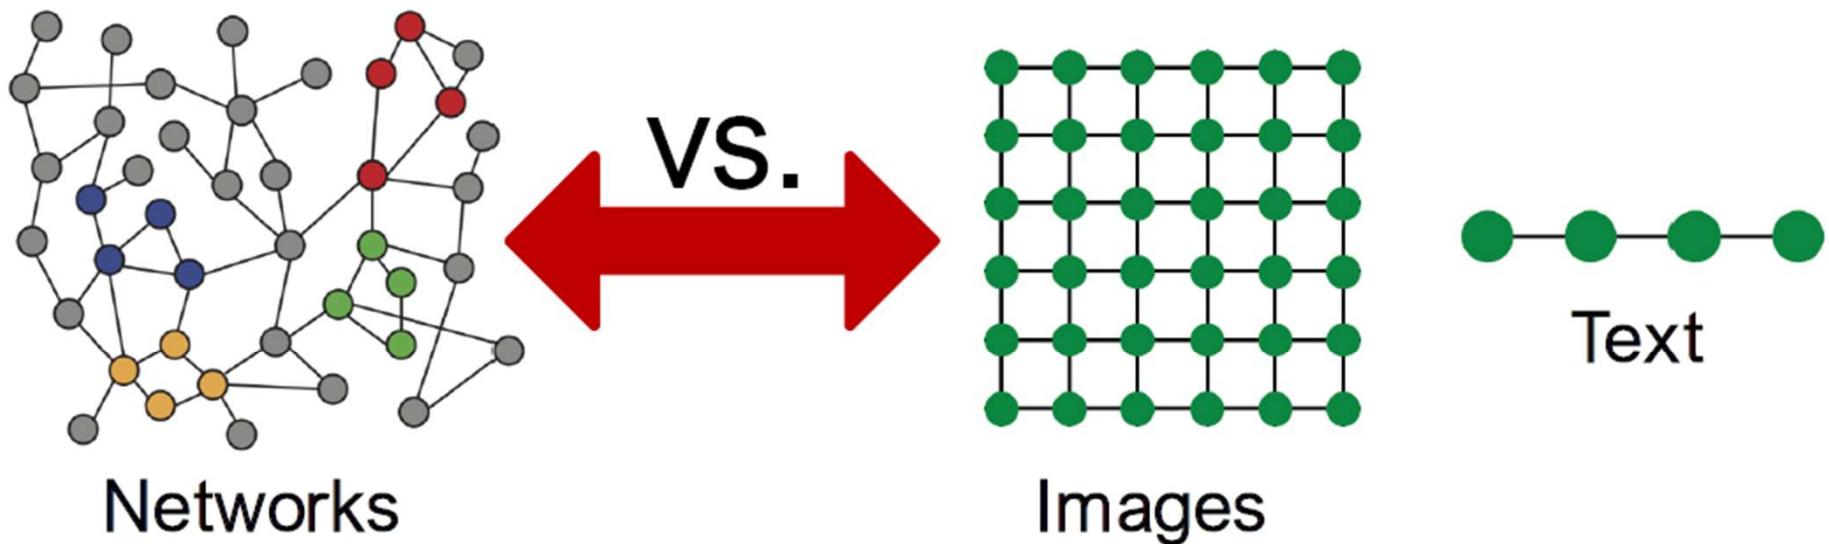
Modern deep learning toolbox is designed  
for simple sequences & grids

# Limitations

- Need new architectures for Graphs

**But networks are far more complex!**

- Arbitrary size and complex topological structure (i.e., no spatial locality like grids)



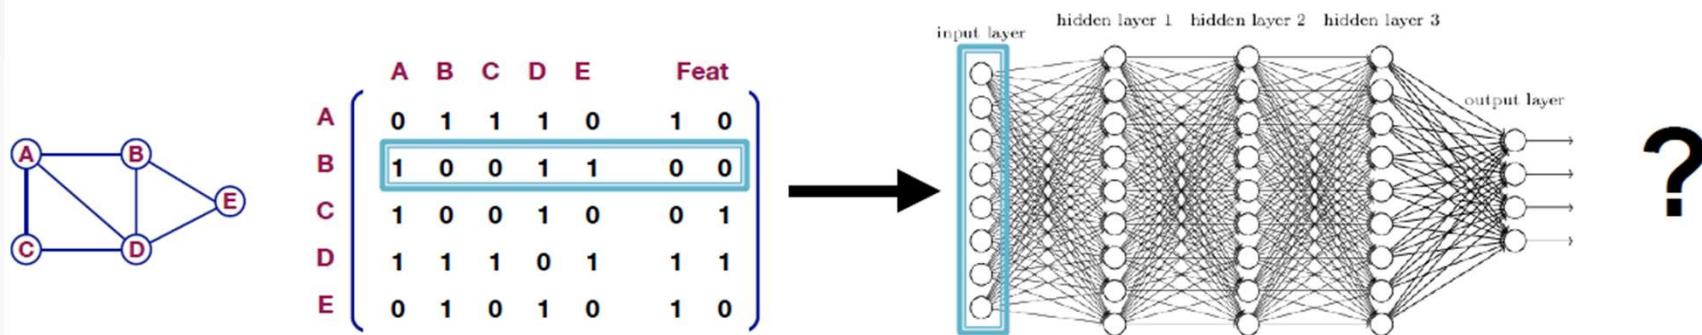
- No fixed node ordering or reference point

# Setup

- Assume we have a graph  $G$ :
  - $V$  is the **vertex set**
  - $A$  is the **adjacency matrix** (assume binary)
  - $X \in \mathbb{R}^{m \times |V|}$  is a matrix of **node features**
  - $v$ : a node in  $V$ ;  $N(v)$ : the set of neighbors of  $v$ .
- **Node features:**
  - Social networks: User profile, User image
  - Biological networks: Gene expression profiles, gene functional information
  - When there is no node feature in the graph dataset:
    - Indicator vectors (one-hot encoding of a node)
    - Vector of constant 1: [1, 1, ..., 1]

# A Naïve Approach

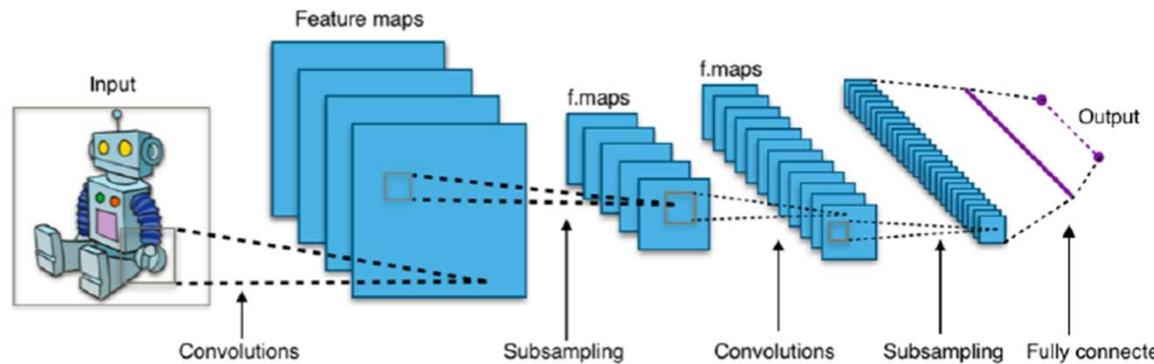
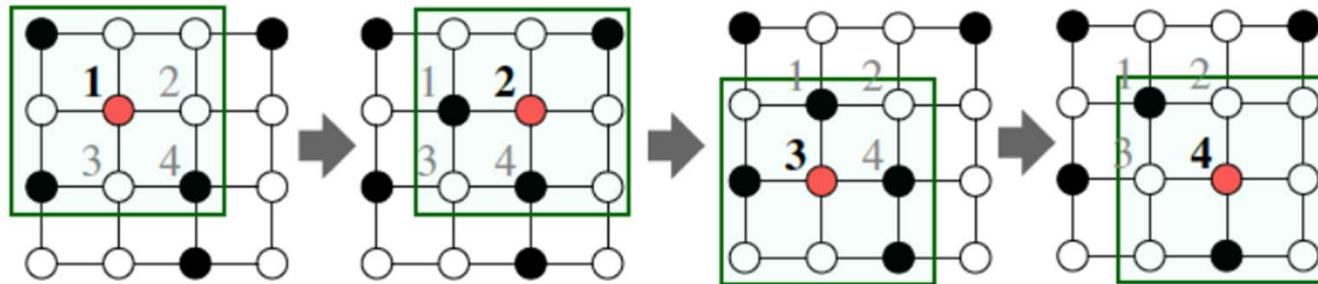
- Join adjacency matrix and features
- Feed them into a deep neural net:



- Issues with this idea:
  - $O(|V|)$  parameters
  - Not applicable to graphs of different sizes
  - Sensitive to node ordering

# Idea: Convolutional Network

## CNN on an image:



Goal is to generalize convolutions beyond simple lattices  
Leverage node features/attributes (e.g., text, images)

# Convolutional Neural Network

- ❖ Inspired by the biology of the visual cortex
  - Local receptive fields are groups of neurons that only respond to a part of what your eyes see (subsampling)
  - They overlap each other to cover the entire visual field (convolutions)
  - They feed into higher layers that identify increasingly complex images
    - Some receptive fields identify horizontal lines, lines at different angles, etc. (filters)
    - These would feed into a layer that identifies shapes
    - Which might feed into a layer that identifies objects
  - For color images, extra layers for red, green, and blue

# How do we know that's a Stop sign?

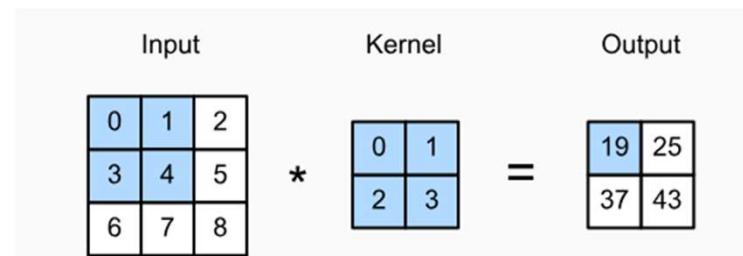
- Individual local receptive fields scan the image looking for edges, and pick up the edges of the stop sign in a layer
- Those edges in turn get picked up by a higher level convolution that identifies the stop sign's shape (and letters, too)
- This shape then gets matched against your pattern of what a stop sign looks like, also using the strong red signal coming from your red layers
- That information keeps getting processed upward until your foot hits the brake!
- A CNN works the same way



# CNN - Components

## ■ Convolution Layer

- This layer is the first layer that is used to extract the various features from the input images. In this layer, the mathematical operation of convolution is performed between the input image and a filter of a particular size MxM. By sliding the filter over the input image, the dot product is taken between the filter and the parts of the input image with respect to the size of the filter (MxM).
- The output is termed as the Feature map which gives us information about the image such as the corners and edges. Later, this feature map is fed to other layers to learn several other features of the input image.



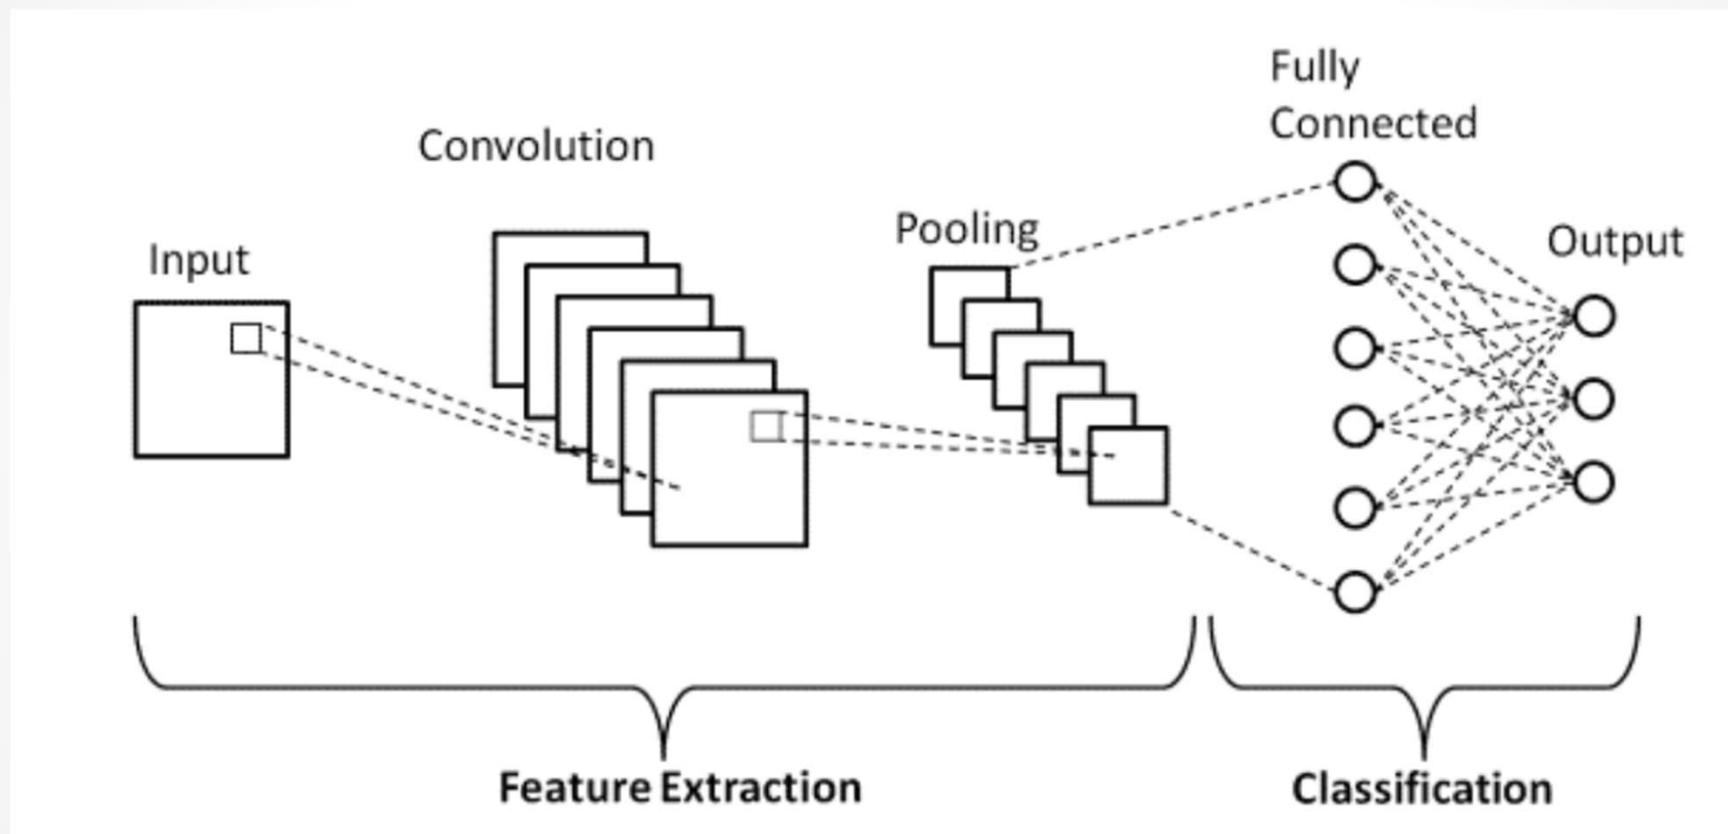
## ■ Pooling Layer

- A Convolutional Layer is followed by a Pooling Layer. The primary aim of this layer is to decrease the size of the convolved feature map to reduce the computational costs. This is performed by decreasing the connections between layers and independently operates on each feature map. Depending upon method used, there are several types of Pooling operations. It basically summarizes the features generated by a convolution layer.
- In Max Pooling, the largest element is taken from feature map. Average Pooling calculates the average of the elements in a predefined sized Image section. The total sum of the elements in the predefined section is computed in Sum Pooling. The Pooling Layer usually serves as a bridge between the Convolutional Layer and the FC Layer.

## • Fully Connected Layer

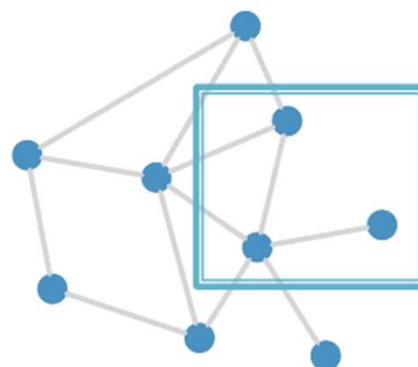
- The Fully Connected (FC) layer consists of the weights and biases along with the neurons and is used to connect the neurons between two different layers. These layers are usually placed before the output layer and form the last few layers of a CNN Architecture.
- In this, the input image from the previous layers are flattened and fed to the FC layer. The flattened vector then undergoes few more FC layers where the mathematical functions operations usually take place. In this stage, the classification process begins to take place.

# CNN - Architecture

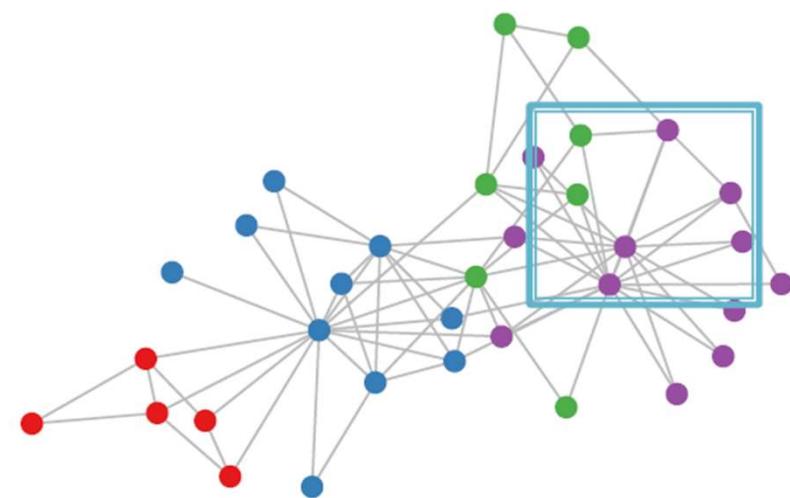


# Real-world Graphs

**But our graphs look like this:**



or this:

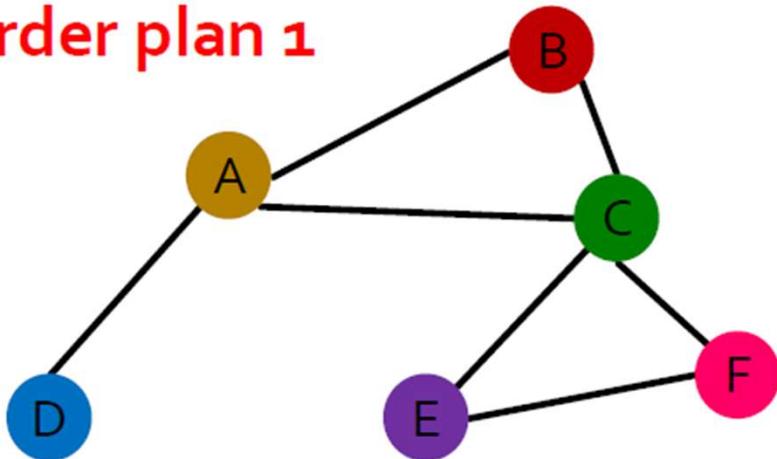


- There is no fixed notion of locality or sliding window on the graph
- Graph is permutation invariant

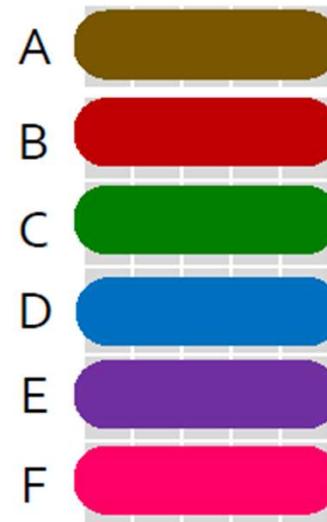
# Permutation Invariance

- Graph does not have a canonical order of the nodes!

Order plan 1



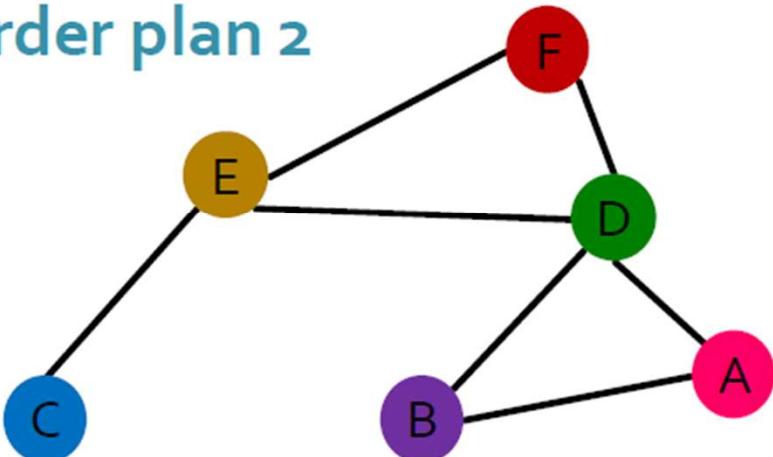
Node features  $X_1$



Adjacency matrix  $A_1$

	A	B	C	D	E	F
A	1	0	1	1	0	0
B	0	1	1	0	0	0
C	1	1	0	0	0	0
D	1	0	0	0	0	0
E	0	0	0	0	0	0
F	0	0	0	0	0	0

Order plan 2



Node features  $X_2$



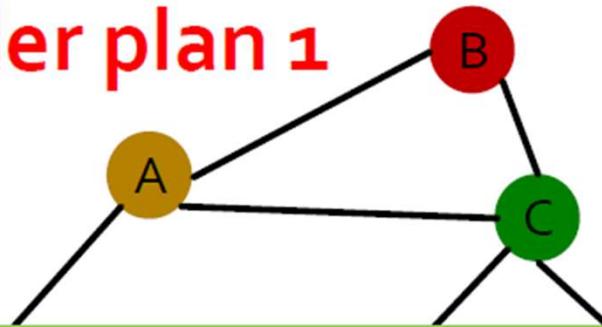
Adjacency matrix  $A_2$

	A	B	C	D	E	F
A	1	0	0	0	0	0
B	0	1	0	0	0	0
C	0	0	1	1	0	0
D	0	0	1	0	0	0
E	0	0	0	0	0	0
F	0	0	0	0	0	0

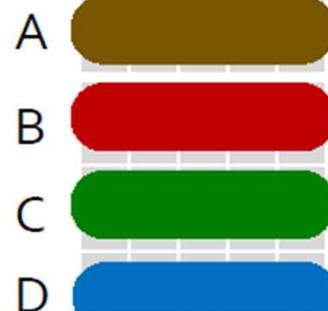
# Permutation Invariance

- Graph does not have a canonical order of the nodes!

Order plan 1



Node feature  $X_1$

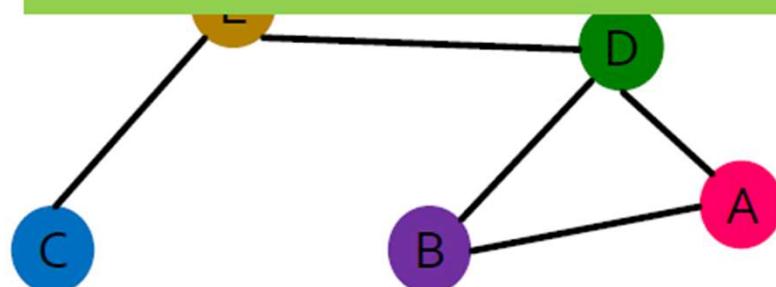


Adjacency matrix  $A_1$

	A	B	C	D	E	F
A	Gray	Blue	Gray	Gray	Gray	Gray
B	Blue	Gray	Gray	Blue	Gray	Gray
C	Gray	Gray	Gray	Gray	Blue	Gray
D	Gray	Gray	Gray	Gray	Gray	Gray

Graph and node representations  
should be the same for Order plan 1  
and Order plan 2

Or



	B	C	D	E	F
B	Gray	Gray	Gray	Blue	Gray
C	Gray	Gray	Gray	Gray	Gray
D	Gray	Gray	Gray	Gray	Gray
E	Gray	Gray	Gray	Gray	Gray
F	Gray	Gray	Gray	Gray	Gray

# Permutation Invariance

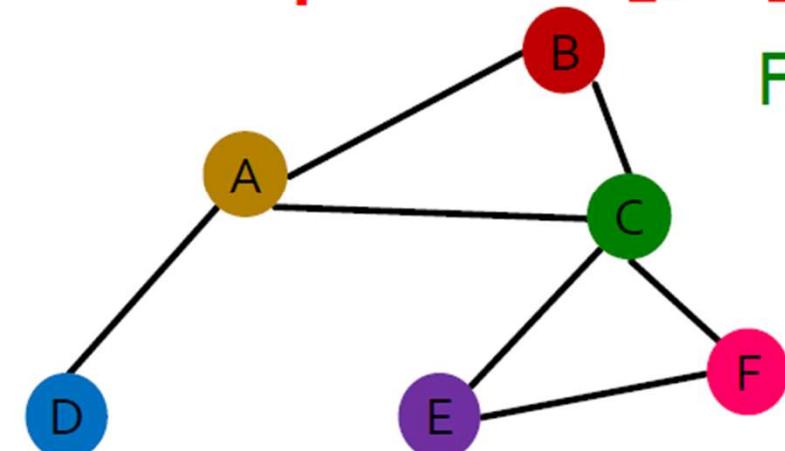
**What does it mean by “graph representation is same for two order plans”?**

- Consider we learn a function  $f$  that maps a graph  $G = (A, X)$  to a vector  $\mathbb{R}^d$  then

$$f(A_1, X_1) = f(A_2, X_2)$$

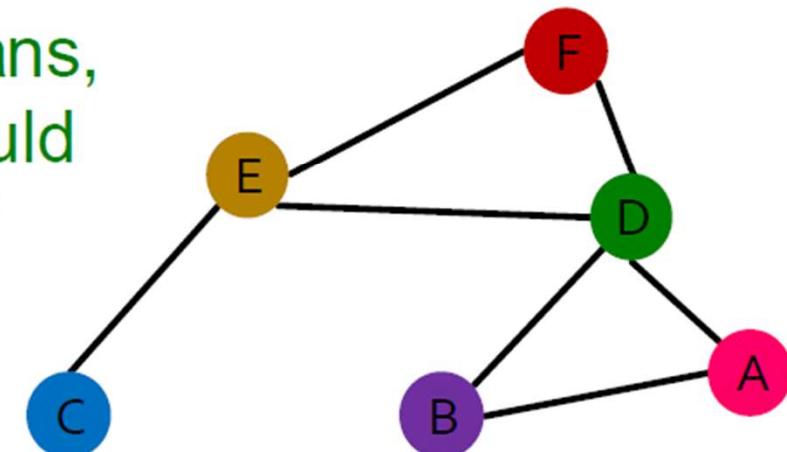
$A$  is the adjacency matrix  
 $X$  is the node feature matrix

**Order plan 1:  $A_1, X_1$**



For two order plans,  
output of  $f$  should  
be the same!

**Order plan 2:  $A_2, X_2$**



# Permutation Invariance

**What does it mean by “graph representation is same for two order plans”?**

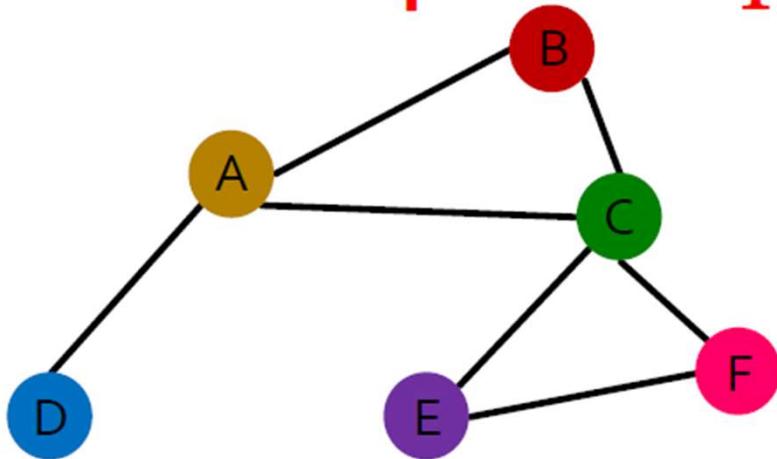
- Consider we learn a function  $f$  that maps a graph  $G = (A, X)$  to a vector  $\mathbb{R}^d$ .  
 $A$  is the adjacency matrix  
 $X$  is the node feature matrix
- Then, if  $f(A_i, X_i) = f(A_j, X_j)$  for any order plan  $i$  and  $j$ , we formally say  $f$  is a **permutation invariant function**.  
For a graph with  $|V|$  nodes, there are  $|V|!$  different order plans.
- Definition:** For any **graph** function  $f: \mathbb{R}^{|V| \times m} \times \mathbb{R}^{|V| \times |V|} \rightarrow \mathbb{R}^d$ ,  $f$  is **permutation-invariant** if  $f(A, X) = f(PAP^T, PX)$  for any permutation  $P$ .

Permutation  $P$ : a shuffle of the node order  
Example: (A,B,C)->(B,C,A)

# Permutation Equivariance

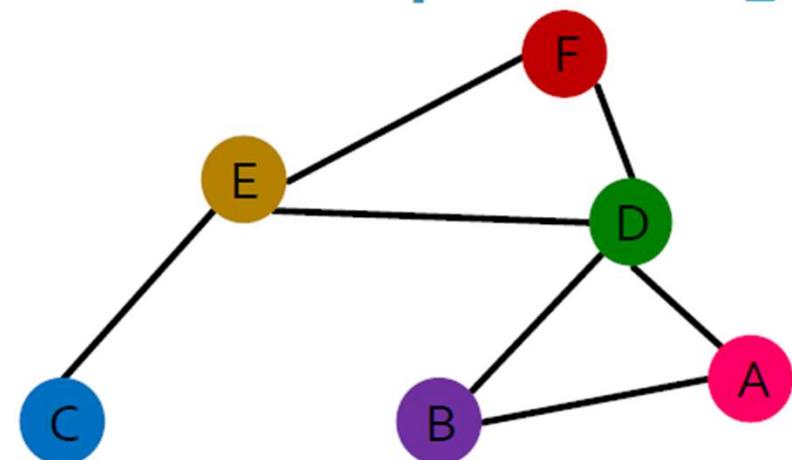
**For node representation:** We learn a function  $f$  that maps nodes of  $G$  to a matrix  $\mathbb{R}^{m \times d}$ .

Order plan 1:  $A_1, X_1$



$$f(A_1, X_1) = \begin{matrix} & \text{A} & \text{B} & \text{C} & \text{D} & \text{E} & \text{F} \\ \text{A} & \text{Yellow} & \text{Red} & \text{Green} & \text{Blue} & \text{Purple} & \text{Red} \\ \text{B} & \text{Red} & \text{Yellow} & \text{Red} & \text{Red} & \text{Red} & \text{Red} \\ \text{C} & \text{Green} & \text{Red} & \text{Yellow} & \text{Red} & \text{Red} & \text{Red} \\ \text{D} & \text{Blue} & \text{Red} & \text{Red} & \text{Yellow} & \text{Red} & \text{Red} \\ \text{E} & \text{Purple} & \text{Red} & \text{Red} & \text{Red} & \text{Yellow} & \text{Red} \\ \text{F} & \text{Red} & \text{Red} & \text{Red} & \text{Red} & \text{Red} & \text{Yellow} \end{matrix}$$

Order plan 2:  $A_2, X_2$

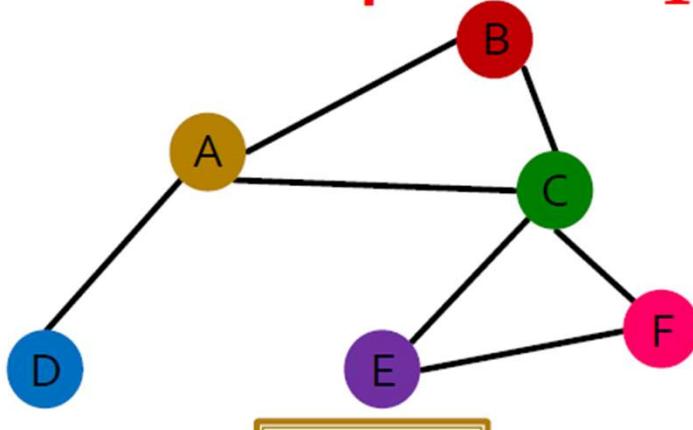


$$f(A_2, X_2) = \begin{matrix} & \text{A} & \text{B} & \text{C} & \text{D} & \text{E} & \text{F} \\ \text{A} & \text{Red} & \text{Red} & \text{Red} & \text{Red} & \text{Yellow} & \text{Red} \\ \text{B} & \text{Purple} & \text{Red} & \text{Red} & \text{Red} & \text{Yellow} & \text{Red} \\ \text{C} & \text{Blue} & \text{Red} & \text{Red} & \text{Red} & \text{Yellow} & \text{Red} \\ \text{D} & \text{Green} & \text{Red} & \text{Red} & \text{Red} & \text{Yellow} & \text{Red} \\ \text{E} & \text{Yellow} & \text{Red} & \text{Red} & \text{Red} & \text{Yellow} & \text{Red} \\ \text{F} & \text{Red} & \text{Red} & \text{Red} & \text{Red} & \text{Yellow} & \text{Red} \end{matrix}$$

# Permutation Equivariance

**For node representation:** We learn a function  $f$  that maps nodes of  $G$  to a matrix  $\mathbb{R}^{m \times d}$ .

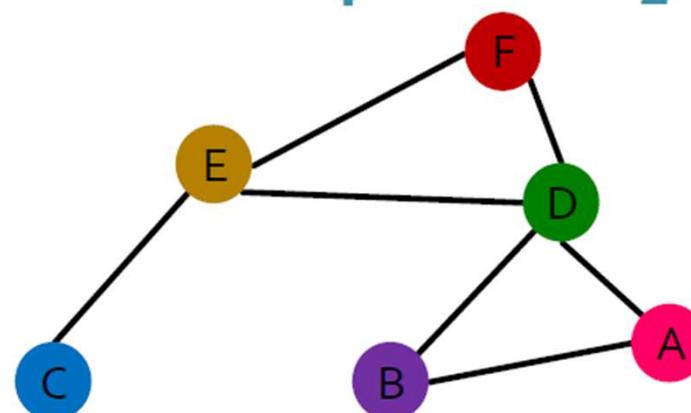
Order plan 1:  $A_1, X_1$



Representation vector  
of the brown node A

$$f(A_1, X_1) = \begin{matrix} & \boxed{\begin{matrix} A & \text{brown} \\ B & \text{red} \\ C & \text{green} \\ D & \text{blue} \\ E & \text{purple} \\ F & \text{pink} \end{matrix}} \\ \begin{matrix} A \\ B \\ C \\ D \\ E \\ F \end{matrix} & \left[ \begin{matrix} \text{brown} & \text{brown} & \text{red} & \text{green} & \text{blue} & \text{purple} \\ \text{red} & \text{brown} & \text{red} & \text{green} & \text{blue} & \text{purple} \\ \text{green} & \text{red} & \text{brown} & \text{red} & \text{blue} & \text{purple} \\ \text{blue} & \text{green} & \text{red} & \text{brown} & \text{red} & \text{purple} \\ \text{purple} & \text{blue} & \text{green} & \text{red} & \text{brown} & \text{red} \\ \text{pink} & \text{purple} & \text{blue} & \text{green} & \text{red} & \text{brown} \end{matrix} \right] \end{matrix}$$

Order plan 2:  $A_2, X_2$



$$f(A_2, X_2) =$$

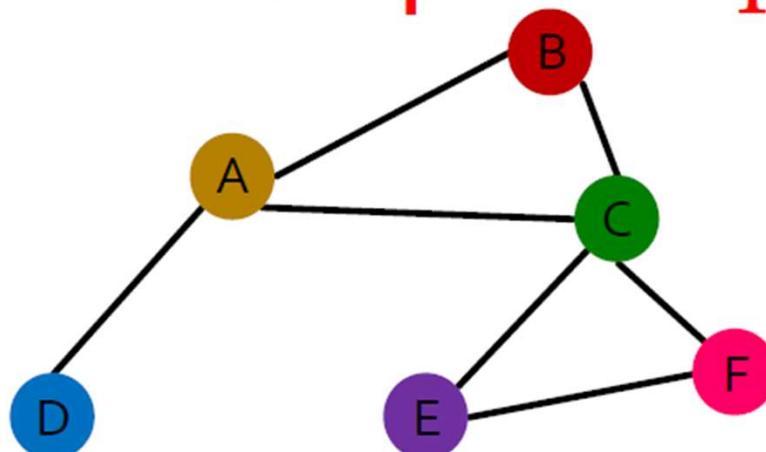
$$\begin{matrix} & \boxed{\begin{matrix} A & \text{red} \\ B & \text{purple} \\ C & \text{blue} \\ D & \text{green} \\ E & \text{brown} \\ F & \text{red} \end{matrix}} \\ \begin{matrix} A \\ B \\ C \\ D \\ E \\ F \end{matrix} & \left[ \begin{matrix} \text{red} & \text{red} & \text{purple} & \text{blue} & \text{green} & \text{brown} \\ \text{purple} & \text{red} & \text{purple} & \text{blue} & \text{green} & \text{brown} \\ \text{blue} & \text{purple} & \text{red} & \text{brown} & \text{green} & \text{brown} \\ \text{green} & \text{blue} & \text{purple} & \text{red} & \text{brown} & \text{brown} \\ \text{brown} & \text{green} & \text{blue} & \text{purple} & \text{red} & \text{red} \\ \text{red} & \text{brown} & \text{green} & \text{blue} & \text{purple} & \text{brown} \end{matrix} \right] \end{matrix}$$

For two order plans, the vector of node at  
the same position in the graph is the same!

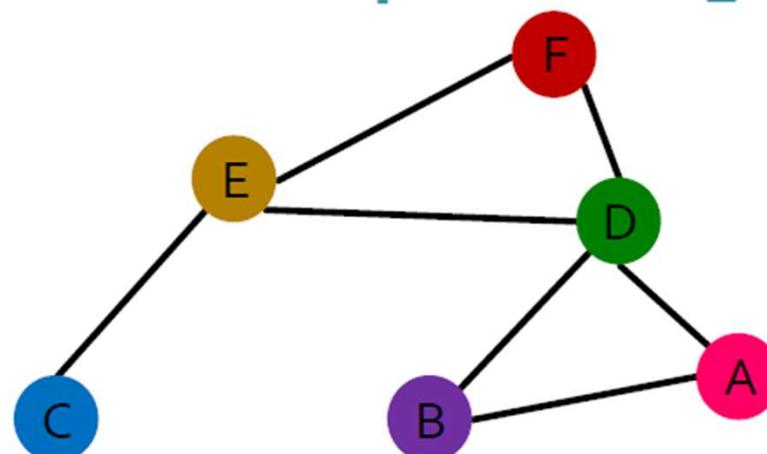
# Permutation Equivariance

**For node representation:** We learn a function  $f$  that maps nodes of  $G$  to a matrix  $\mathbb{R}^{m \times d}$ .

Order plan 1:  $A_1, X_1$



Order plan 2:  $A_2, X_2$



A	Yellow	Yellow
B	Red	Red
C	Green	Green
D	Blue	Blue
E	Purple	Purple
F	Red	Red

Representation vector  
of the green node C

$$f(A_1, X_1) =$$

For two order plans, the vector of node at  
the same position in the graph is the same! F

A	Red	Red
B	Purple	Purple
C	Blue	Blue
D	Green	Green
E	Yellow	Yellow
F	Red	Red

Representation vector  
of the green node D

$$f(A_2, X_2) =$$

# Permutation Equivariance

## For node representation

- Consider we learn a function  $f$  that maps a graph  $G = (A, X)$  to a matrix  $\mathbb{R}^{m \times d}$
- If the output vector of a node at the same position in the graph remains unchanged for any order plan, we say  $f$  is **permutation equivariant**.
- **Definition:** For any **node** function  $f: \mathbb{R}^{|V| \times m} \times \mathbb{R}^{|V| \times |V|} \rightarrow \mathbb{R}^{|V| \times m}$ ,  $f$  is **permutation-equivariant** if  $Pf(A, X) = f(PAP^T, PX)$  for any permutation  $P$ .

# Invariance & Equivariance

## ■ Permutation-invariant

$$f(A, X) = f(PAP^T, PX)$$

Permute the input, the output stays the same.  
(map a graph to a vector)

## ■ Permutation-equivariant

$$Pf(A, X) = f(PAP^T, PX)$$

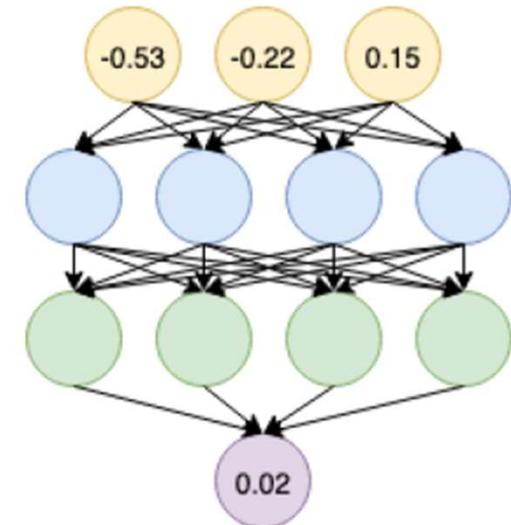
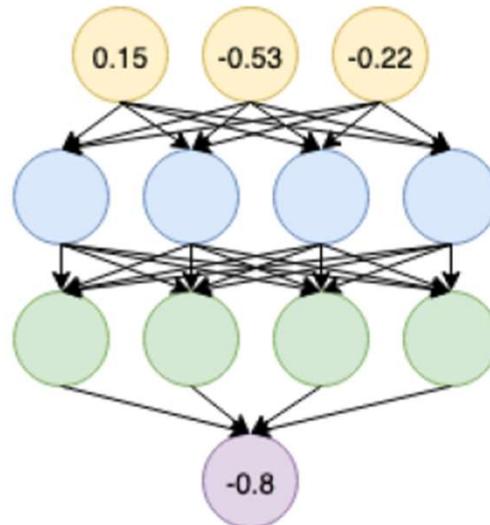
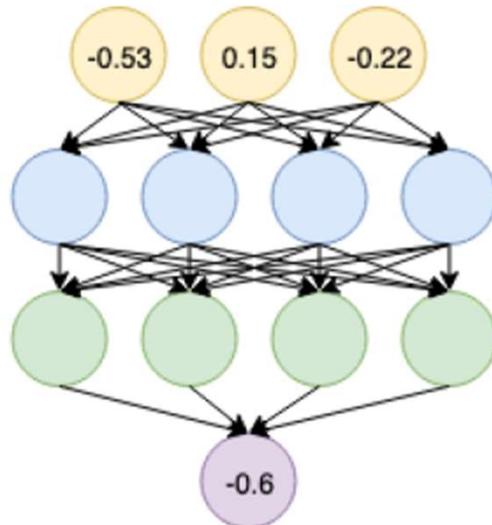
Permute the input, output also permutes accordingly.  
(map a graph to a matrix)

- **Invariance:** Output does not depend on the representation of the input
- **Equivariance:** Output must be permuted when the input is permuted

# Invariance & Equivariance

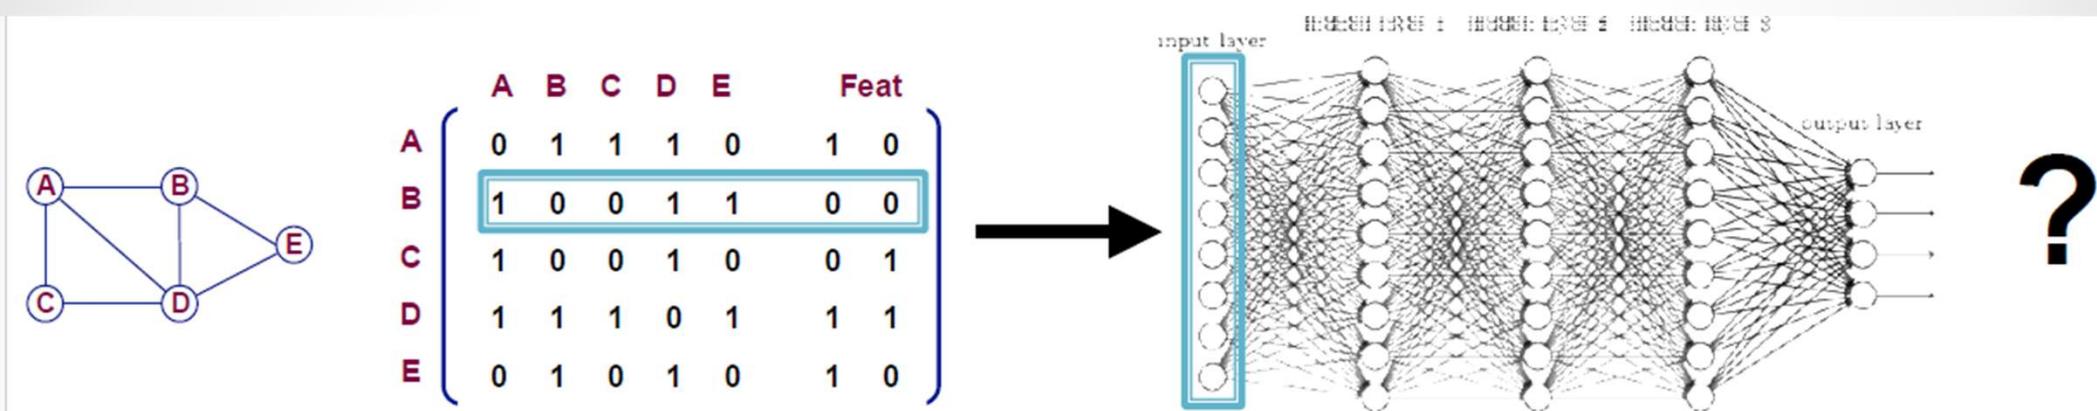
- Are other neural network architectures, e.g., MLPs, permutation invariant / equivariant?
- No.

Switching the order of the input leads to different outputs!



# Invariance & Equivariance

- Are other neural network architectures, e.g., MLPs, permutation invariant / equivariant?
- No.

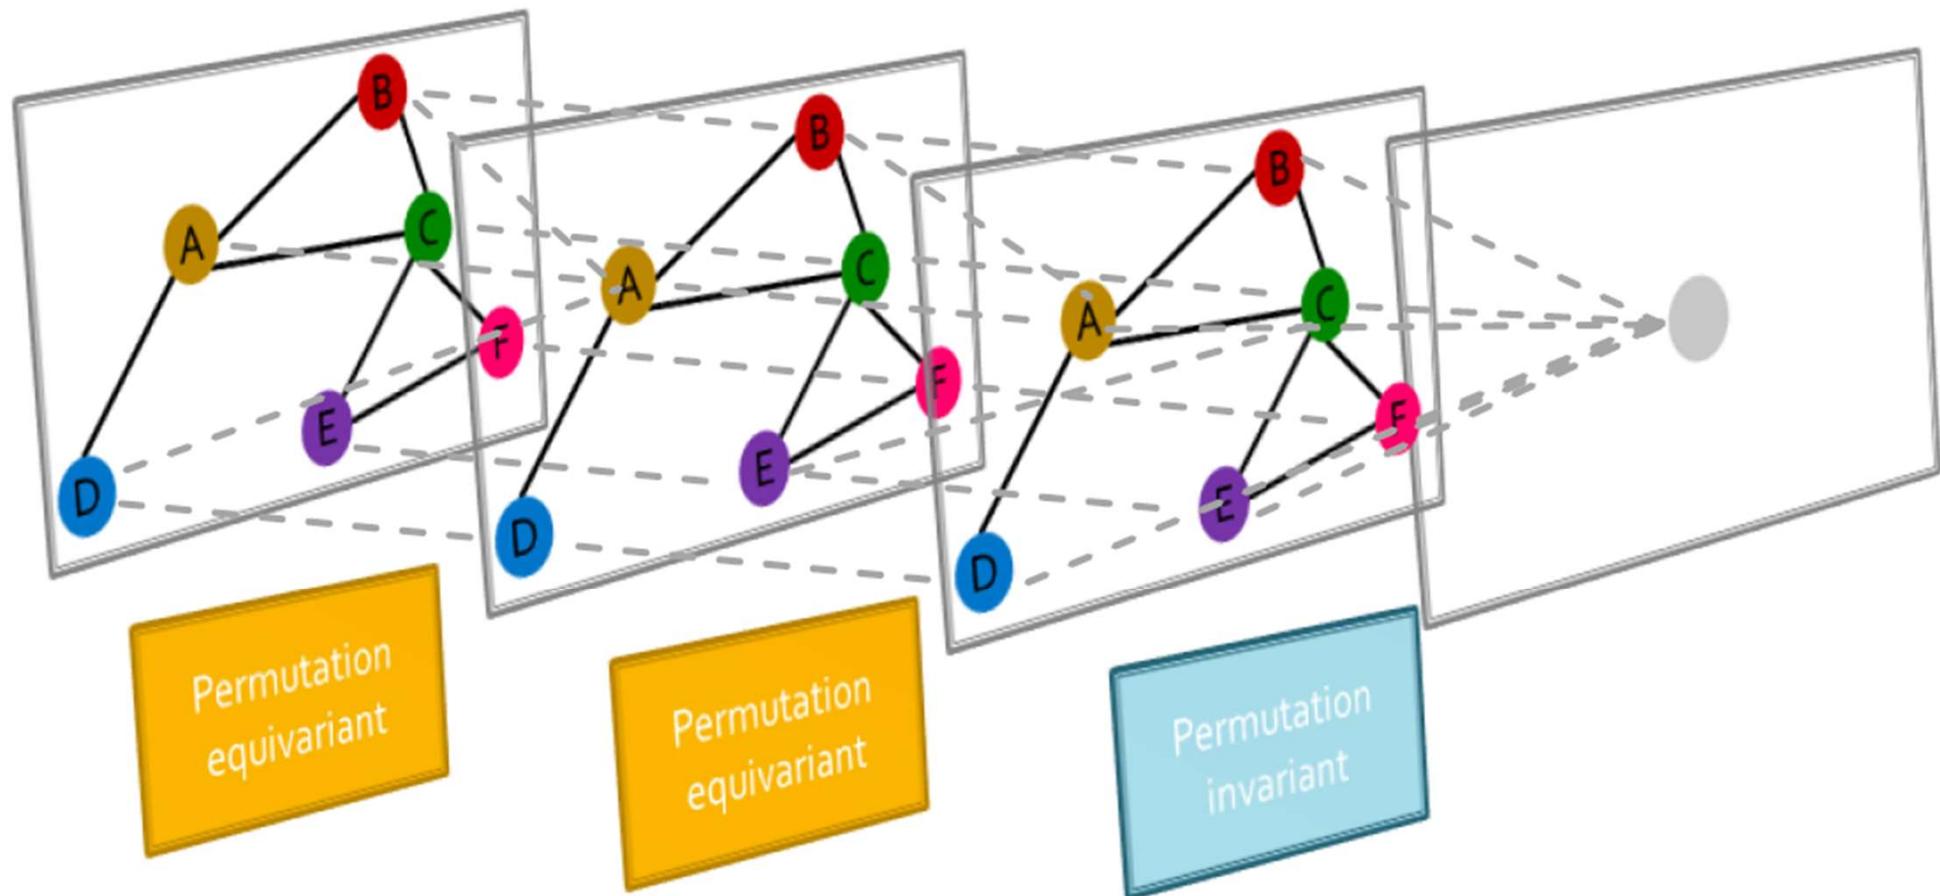


This explains why **the naïve MLP approach fails for graphs!**

Design graph neural networks that are permutation invariant / equivariant

# Graph Neural Networks

- Graph neural networks consist of multiple permutation equivariant / invariant functions.

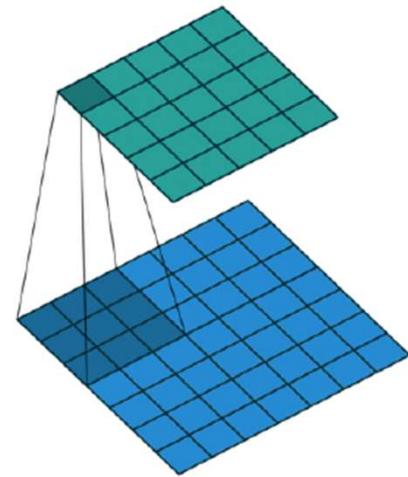


# Overview

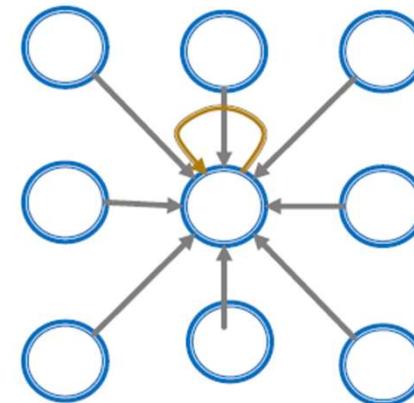
- Shallow Encodings - Limitations
- Deep Learning for Graphs
- Graph Convolution Networks

# From Images to Graphs

Single Convolutional neural network (CNN) layer  
with 3x3 filter:



Image



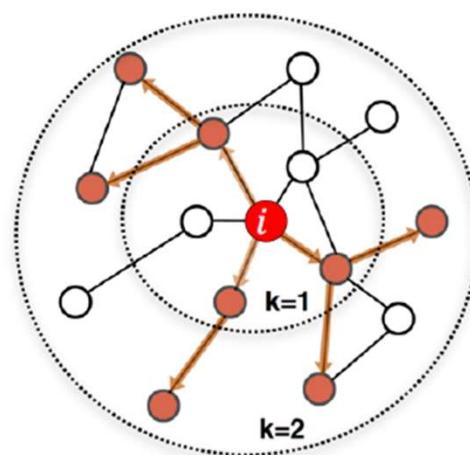
Graph

**Idea:** transform information at the neighbors and combine it:

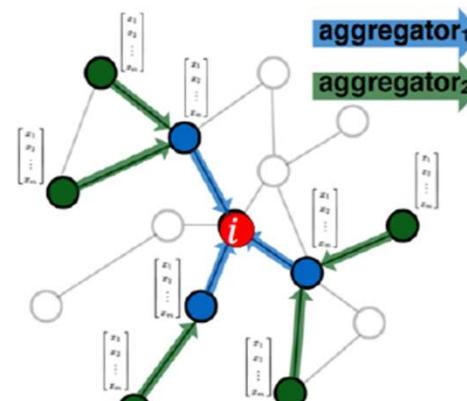
- Transform “messages”  $h_i$  from neighbors:  $W_i h_i$
- Add them up:  $\sum_i W_i h_i$

# Graph Convolutional Network

Idea: Node's neighborhood defines a computation graph



Determine node computation graph



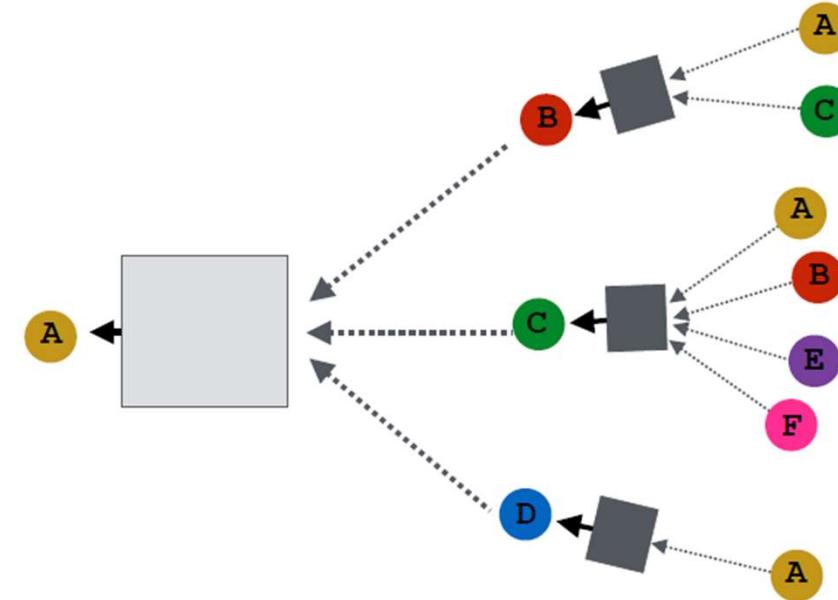
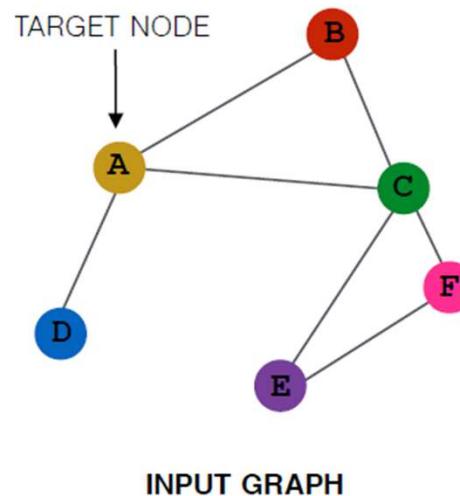
Propagate and transform information

Learn how to propagate information across the graph to compute node features

[Kipf and Welling, ICLR 2017]

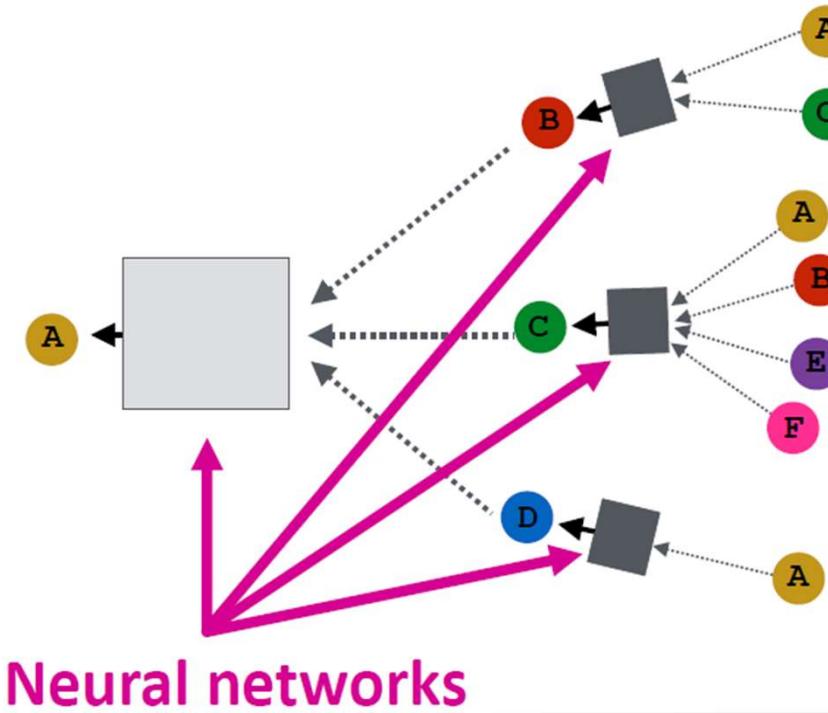
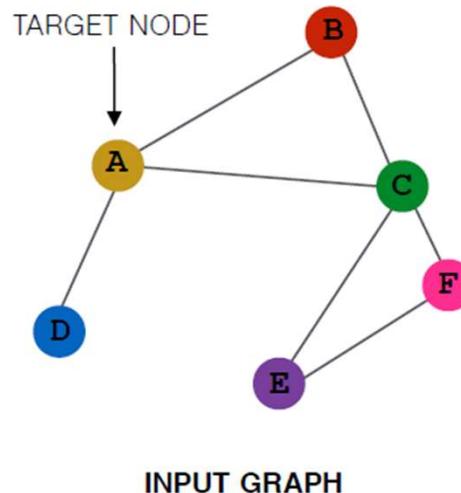
# Idea: Aggregate Neighbors

- Key idea: Generate node embeddings based on local network neighborhoods



# Idea: Aggregate Neighbors

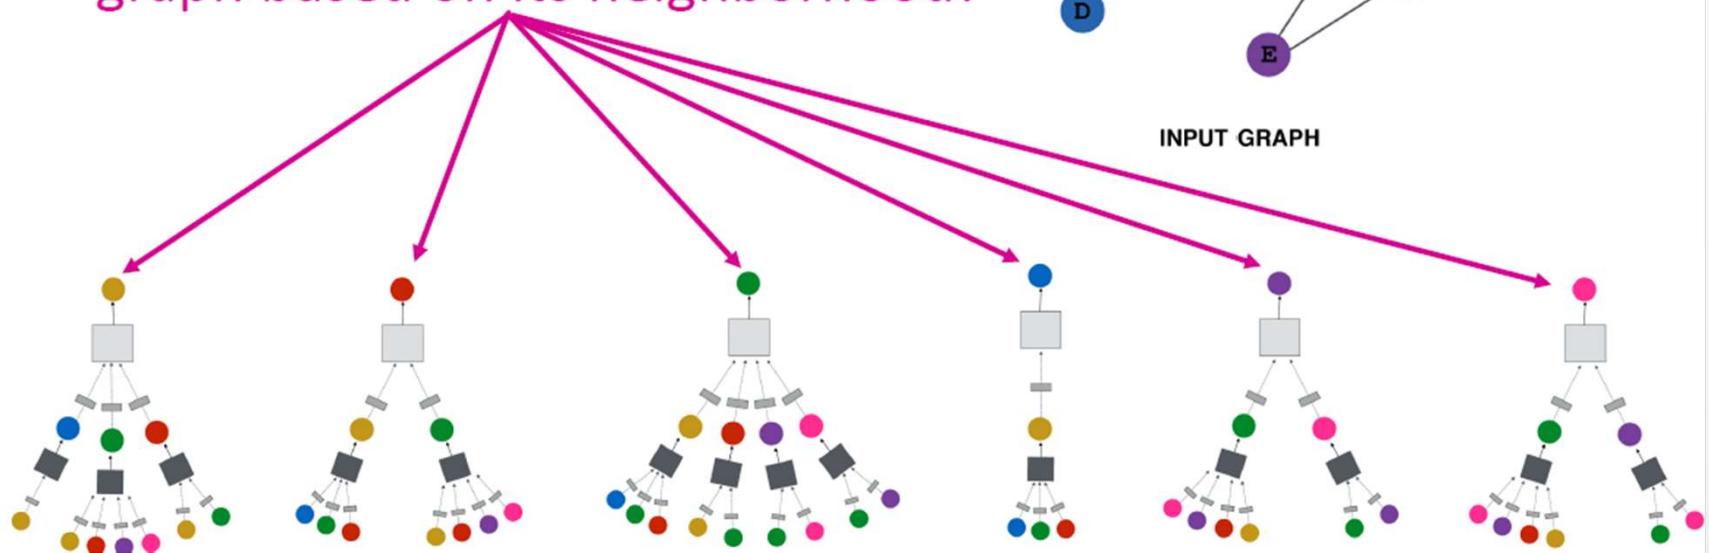
- **Intuition:** Nodes aggregate information from their neighbors using neural networks



# Idea: Aggregate Neighbors

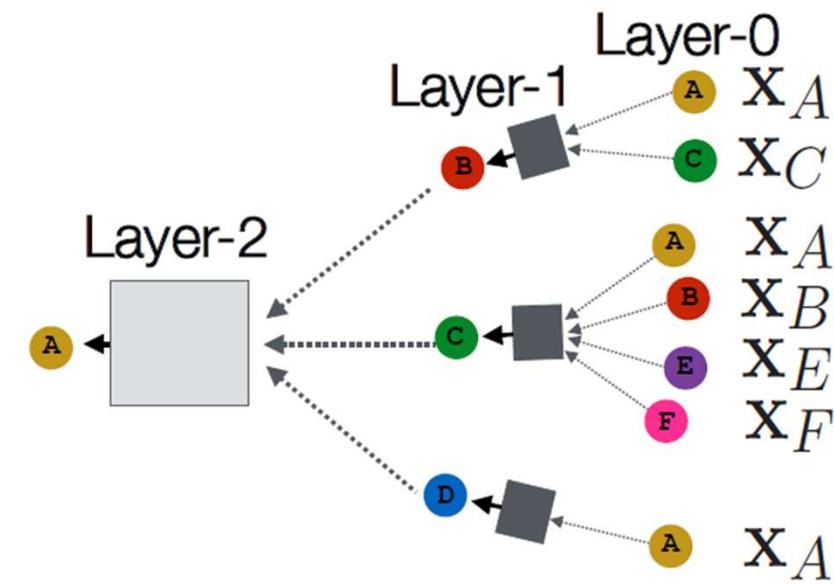
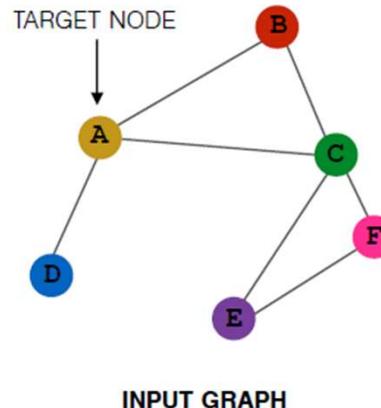
- **Intuition:** Network neighborhood defines a computation graph

Every node defines a computation graph based on its neighborhood!



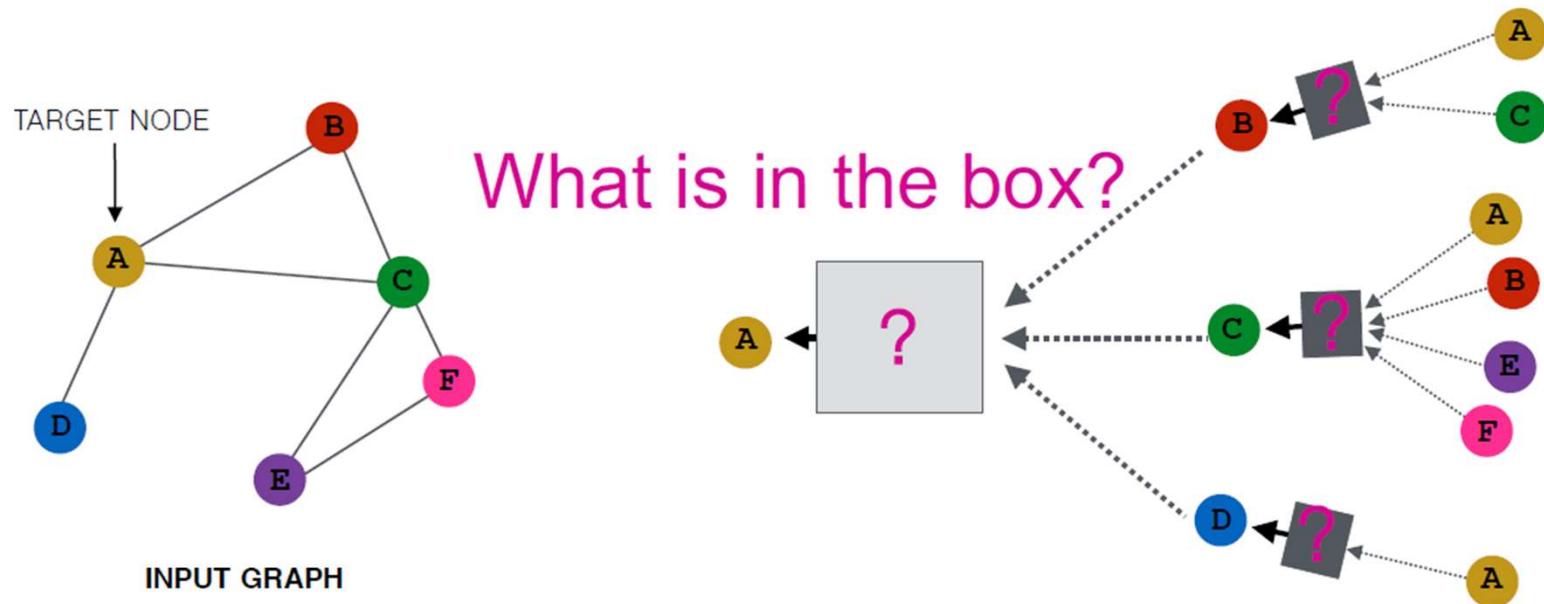
# Deep Model: Many Layers

- Model can be **of arbitrary depth**:
  - Nodes have embeddings at each layer
  - Layer-0 embedding of node  $u$  is its input feature,  $x_u$
  - Layer- $k$  embedding gets information from nodes that are  $K$  hops away



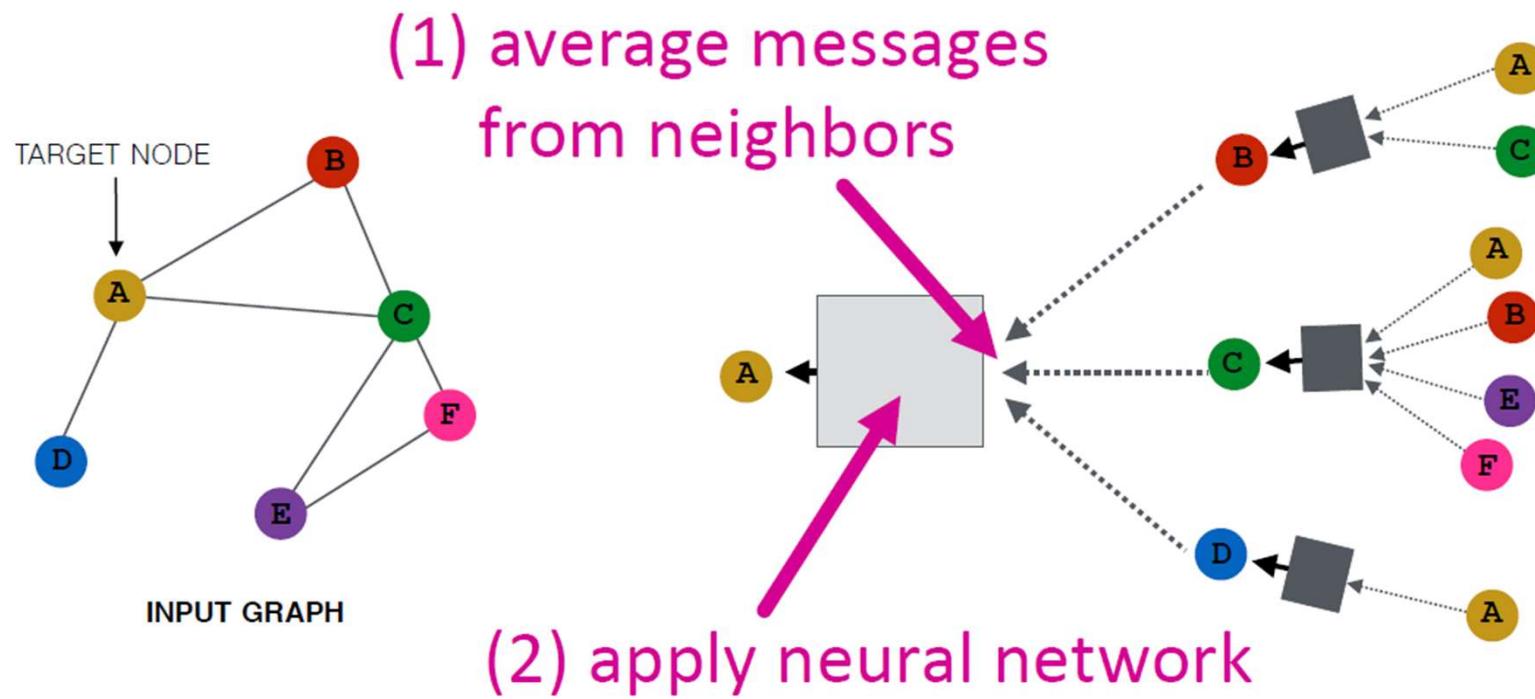
# Neighborhood Aggregation

- **Neighborhood aggregation:** Key distinctions are in how different approaches aggregate information across the layers



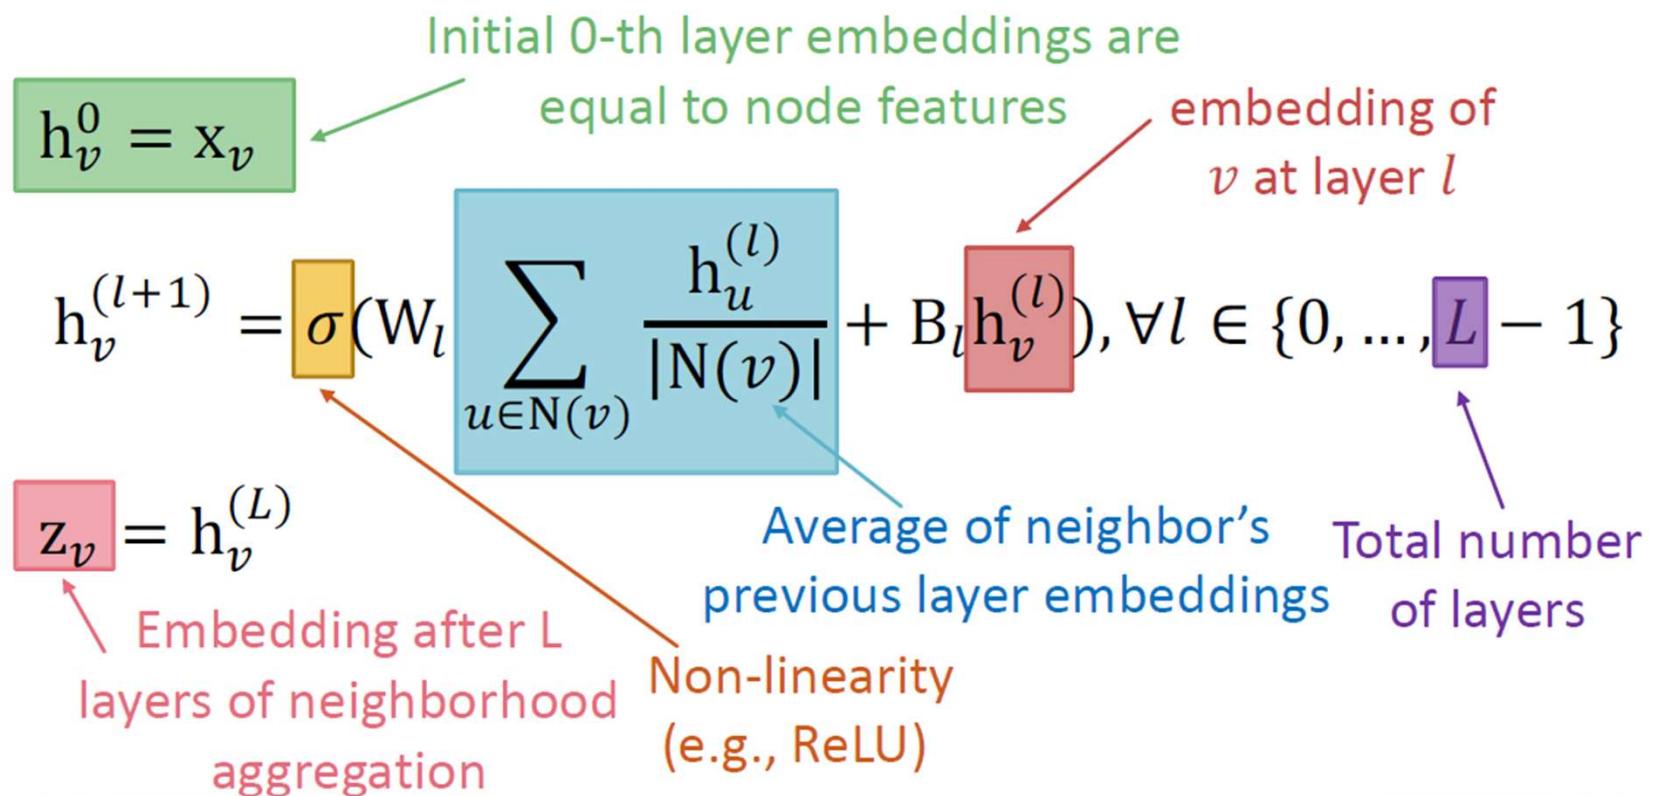
# Neighborhood Aggregation

- **Basic approach:** Average information from neighbors and apply a neural network



# The Math: Deep Encoder

- **Basic approach:** Average neighbor messages and apply a neural network

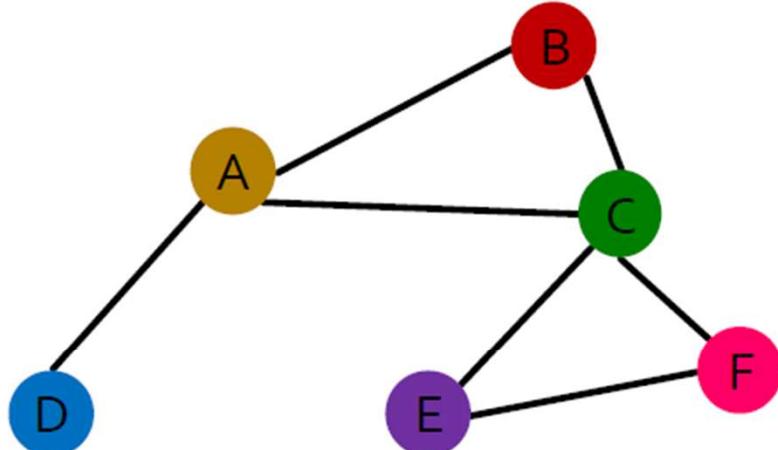


# GCN: Invariance & Equivariance

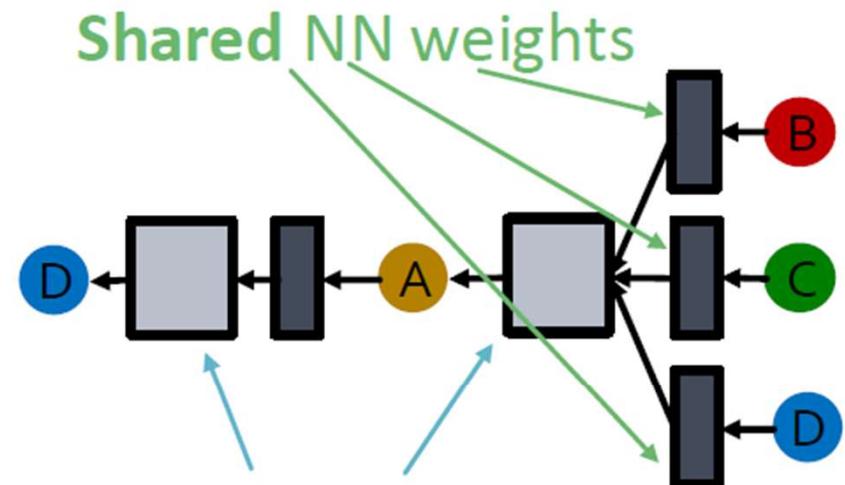
What are the **invariance** and **equivariance** properties for a GCN?

- Given a node, the GCN that computes its embedding is **permutation invariant**

i) Shared NN weights ii) The aggregation is Permutation Invariant



Target Node

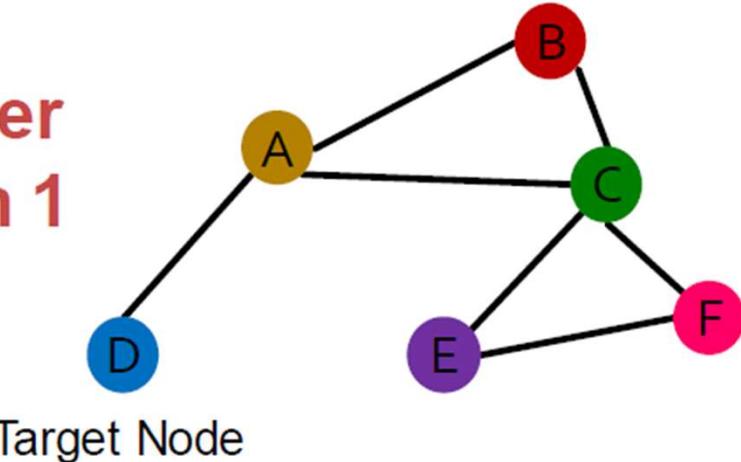


Average of neighbor's previous layer  
embeddings - **Permutation invariant**

# GCN: Invariance & Equivariance

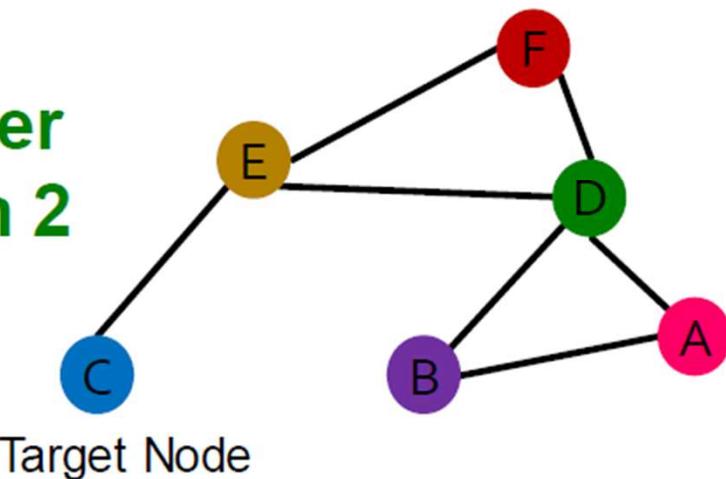
- Considering all nodes in a graph, GCN computation is **permutation equivariant**

Order plan 1



Target Node

Order plan 2



Target Node

Node feature $X_1$
A
B
C
D
E
F

Adjacency matrix $A_1$
A
B
C
D
E
F

Embeddings $H_1$
A
B
C
D
E
F

Permute the input, the output also permutes accordingly - **permutation equivariant**

Node feature $X_2$
A
B
C
D
E
F

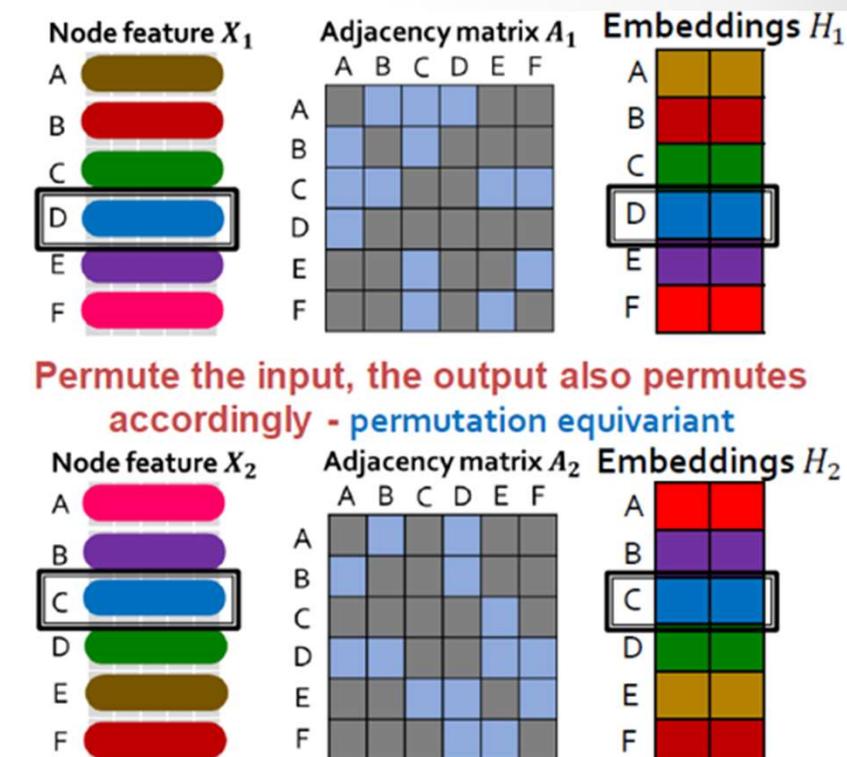
Adjacency matrix $A_2$
A
B
C
D
E
F

Embeddings $H_2$
A
B
C
D
E
F

# GCN: Invariance & Equivariance

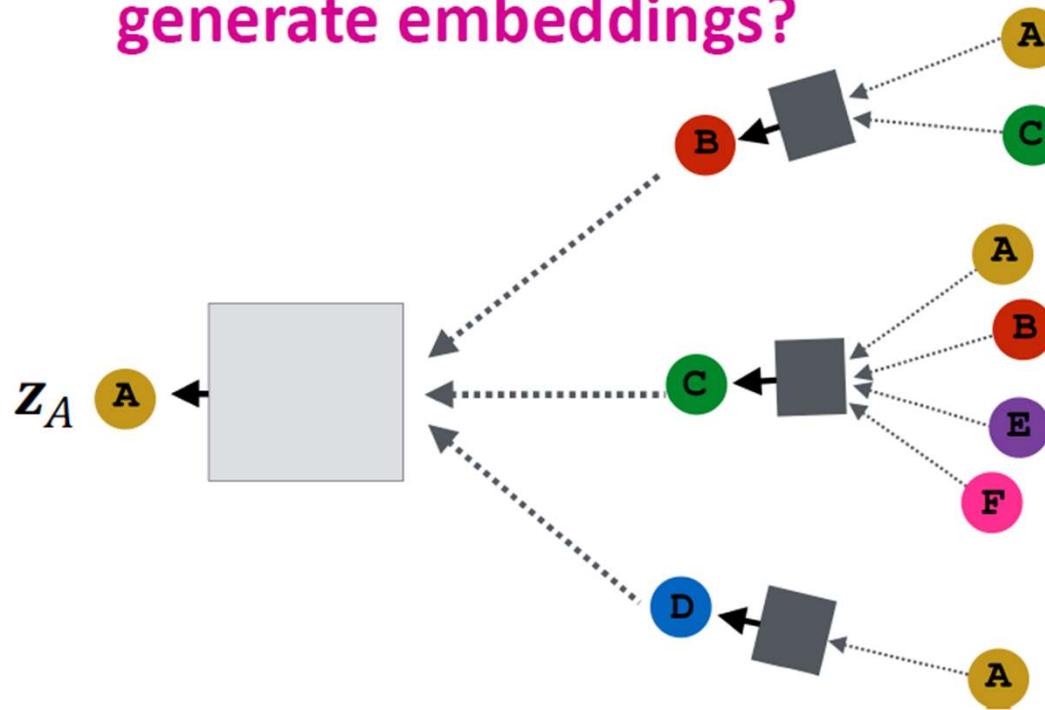
## ■ Considering all nodes in a graph, GCN computation is permutation equivariant

1. The rows of **input node features** and **output embeddings** are aligned
2. We know computing the embedding of **a given node** with GCN is **invariant**.
3. So, after permutation, the **location of a given node** in the **input node feature** matrix is changed, and the **the output embedding of a given changes in the same way**.



# Training the Model

How do we train the model to generate embeddings?



Need to define a loss function on the embeddings

# Model Parameters

$$\begin{aligned} h_v^{(0)} &= x_v \\ h_v^{(l+1)} &= \sigma(W_l \sum_{u \in N(v)} \frac{h_u^{(l)}}{|N(v)|} + B_l h_v^{(l)}), \forall l \in \{0, \dots, L-1\} \\ z_v &= h_v^{(L)} \end{aligned}$$

Trainable weight matrices  
(i.e., what we learn)

Final node embedding

We can feed these **embeddings into any loss function** and run SGD to **train the weight parameters**

$h_v^l$ : the hidden representation of node  $v$  at layer  $l$

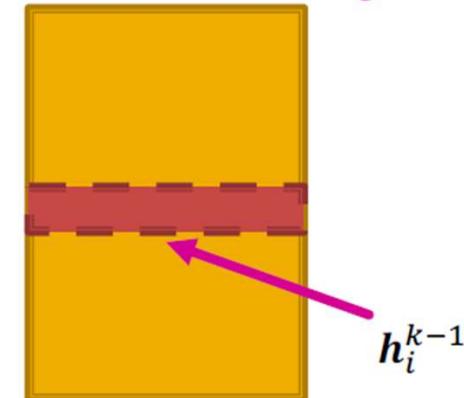
■  $W_k$ : weight matrix for neighborhood aggregation

■  $B_k$ : weight matrix for transforming hidden vector of self

# Matrix Formulation

- Many aggregations can be performed efficiently by (sparse) matrix operations

- Let  $H^{(l)} = [h_1^{(l)} \dots h_{|V|}^{(l)}]^T$  Matrix of hidden embeddings  $H^{k-1}$
- Then:  $\sum_{u \in N_v} h_u^{(l)} = A_{v,:} H^{(l)}$
- Let  $D$  be diagonal matrix where  $D_{v,v} = \text{Deg}(v) = |N(v)|$ 
  - The inverse of  $D$ :  $D^{-1}$  is also diagonal:  
$$D_{v,v}^{-1} = 1/|N(v)|$$
- Therefore,



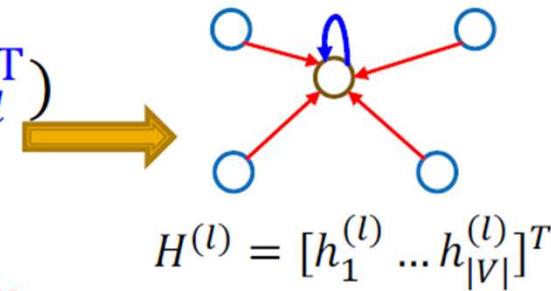
$$\sum_{u \in N(v)} \frac{h_u^{(l-1)}}{|N(v)|} \longrightarrow H^{(l+1)} = D^{-1} A H^{(l)}$$

# Matrix Formulation

- Re-writing update function in matrix form:

$$H^{(l+1)} = \sigma(\tilde{A}H^{(l)}W_l^T + H^{(l)}B_l^T)$$

where  $\tilde{A} = D^{-1}A$



- Red: neighborhood aggregation
- Blue: self transformation
- In practice, this implies that efficient sparse matrix multiplication can be used ( $\tilde{A}$  is sparse)
- **Note:** not all GNNs can be expressed in matrix form, when aggregation function is complex

# How to train a GNN

- Node embedding  $\mathbf{z}_v$  is a function of input graph
- **Supervised setting:** we want to minimize the loss

$$\min_{\Theta} \mathcal{L}(\mathbf{y}, f(\mathbf{z}_v))$$

- $\mathbf{y}$ : node label
- $\mathcal{L}$  could be L2 if  $\mathbf{y}$  is real number, or cross entropy if  $\mathbf{y}$  is categorical
- **Unsupervised setting:**
  - No node label available
  - **Use the graph structure as the supervision!**

# Unsupervised Training

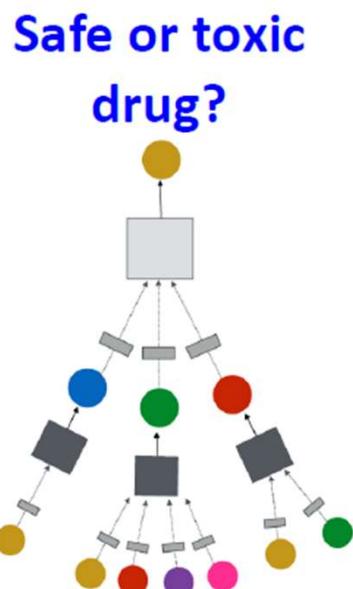
- “Similar” nodes have similar embeddings

$$\mathcal{L} = \sum_{z_u, z_v} \text{CE}(y_{u,v}, \text{DEC}(z_u, z_v))$$

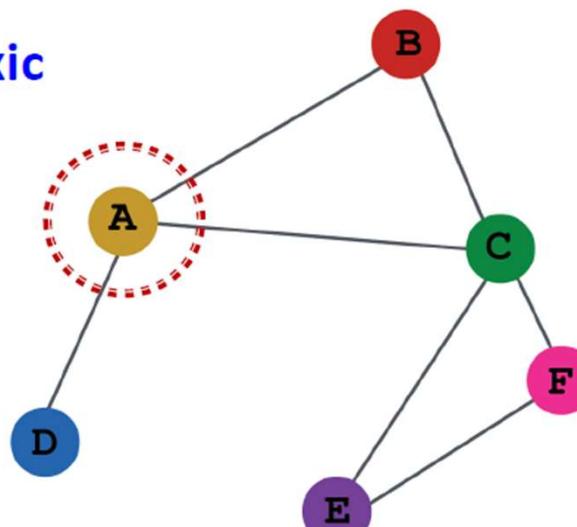
- Where  $y_{u,v} = 1$  when node  $u$  and  $v$  are **similar**
- **CE** is the cross entropy
- **DEC** is the decoder such as inner product
- **Node similarity** a loss based on:
  - **Random walks** (node2vec, DeepWalk)
  - **Matrix factorization**
  - **Node proximity in the graph**

# Supervised Training

**Directly train** the model for a supervised task  
(e.g., node classification)



Safe or toxic drug?

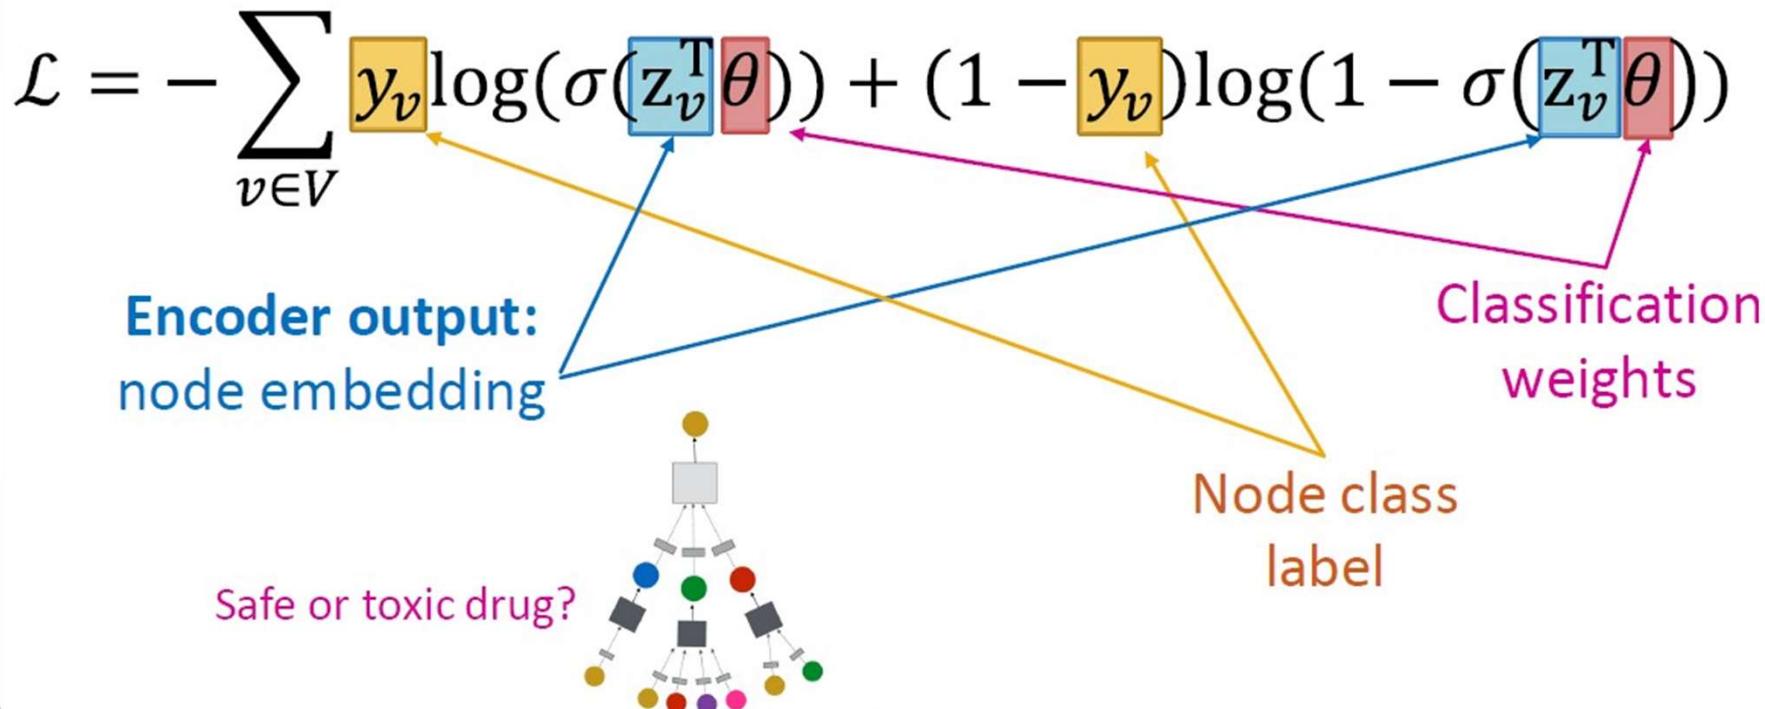


E.g., a drug-drug interaction network

# Supervised Training

**Directly train** the model for a supervised task  
(e.g., **node classification**)

- Use cross entropy loss

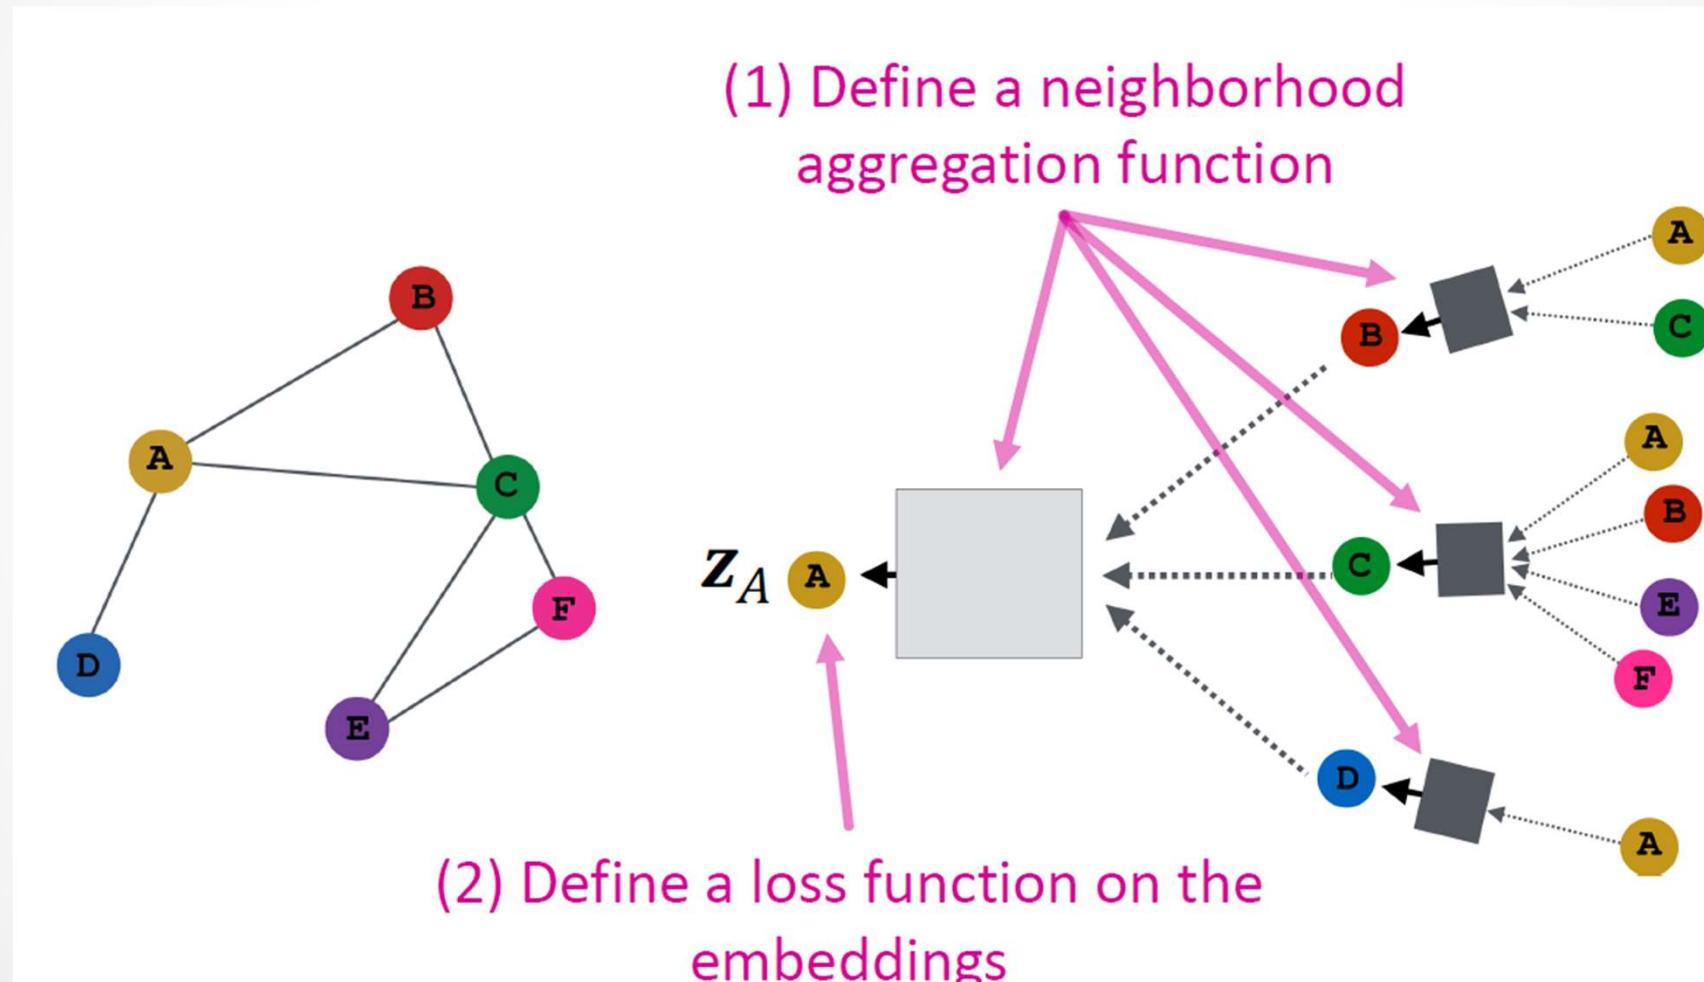


We need to minimize the loss.

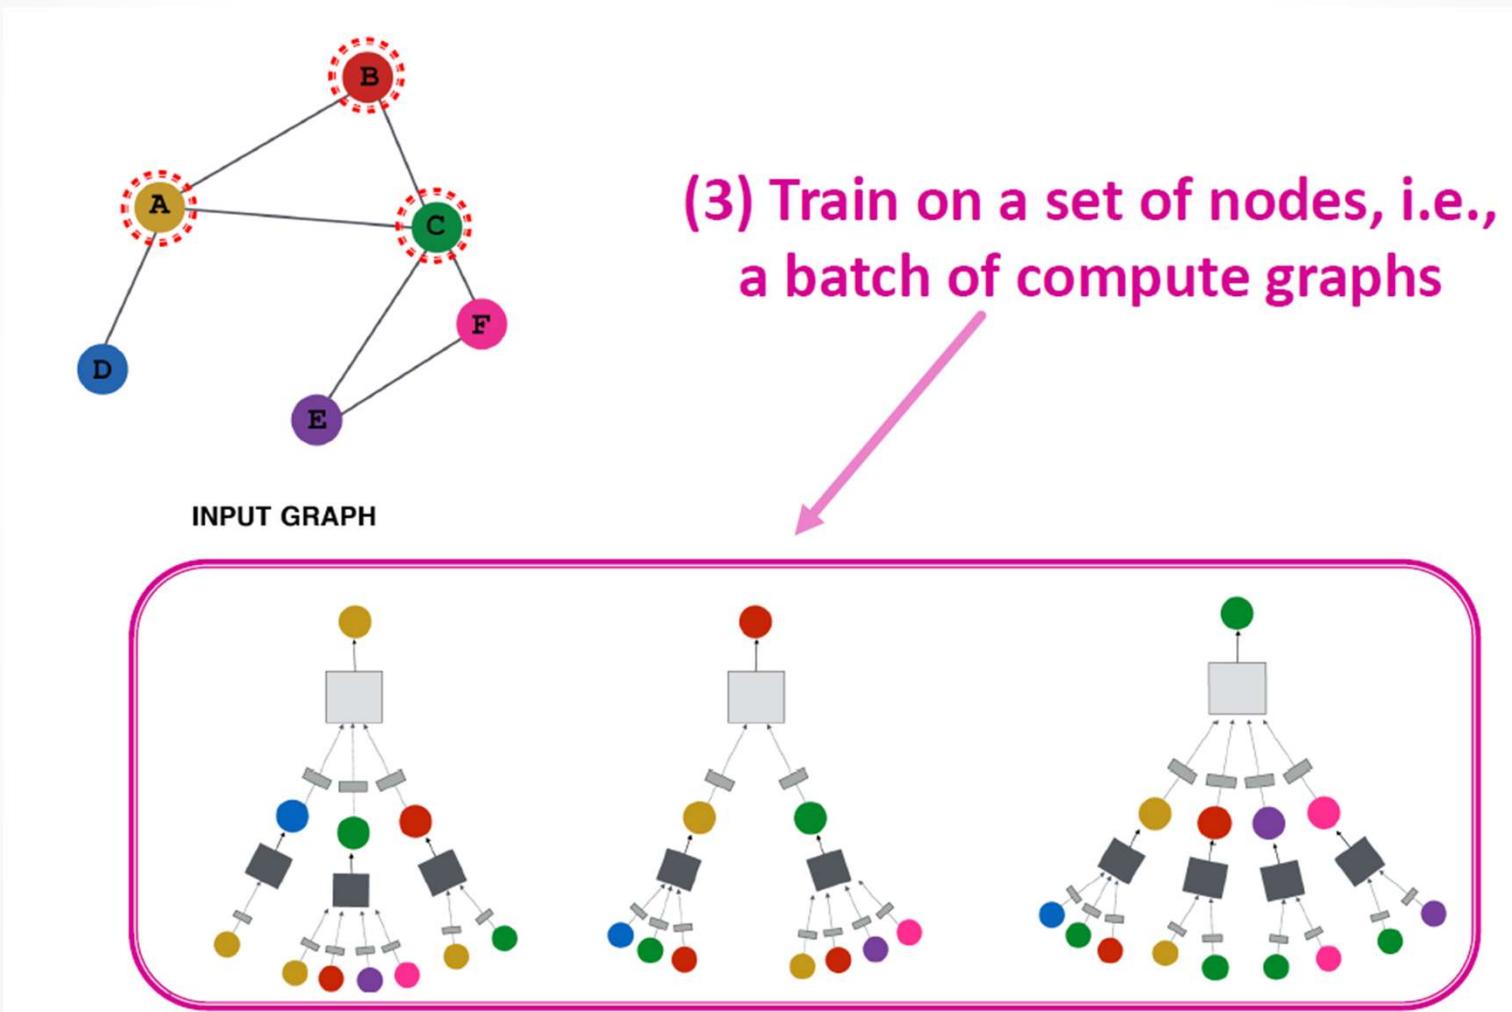
If class label = 1, 1<sup>st</sup> term needs to be high => encoder output needs to be high  
2<sup>nd</sup> term is 0

If class label = 0, 2<sup>nd</sup> term needs to be high => encoder output needs to be low  
1<sup>st</sup> term is 0

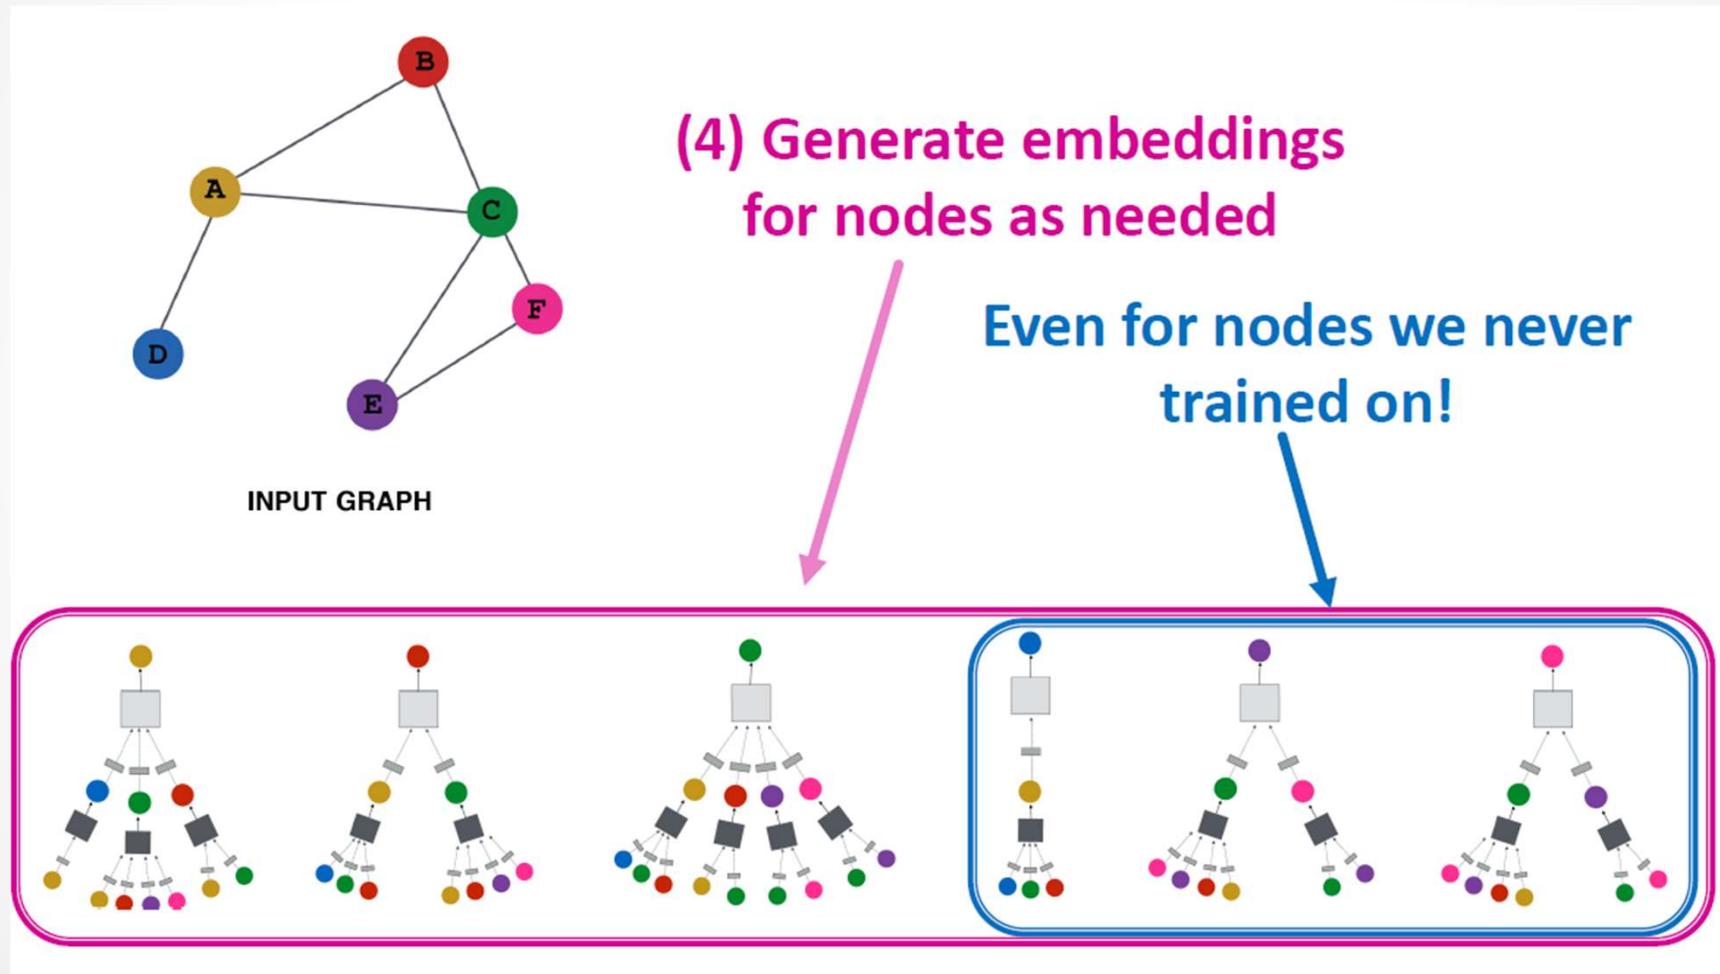
# Model Design: Overview



# Model Design: Overview

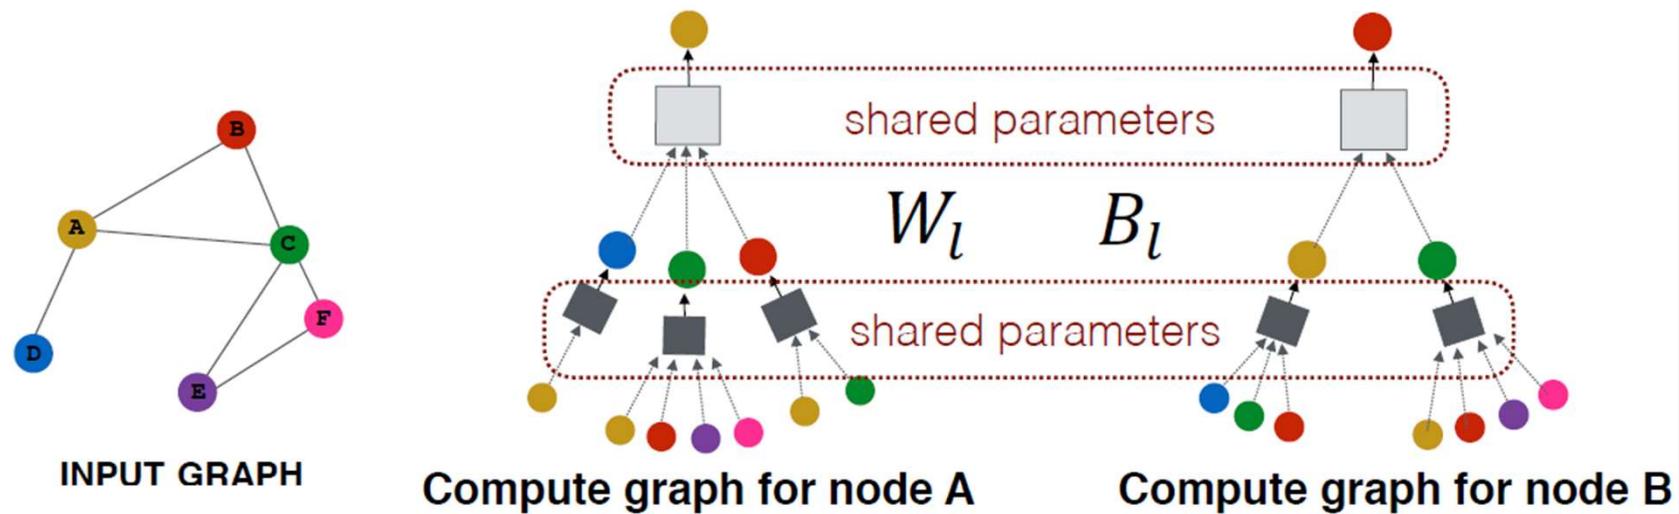


# Model Design: Overview

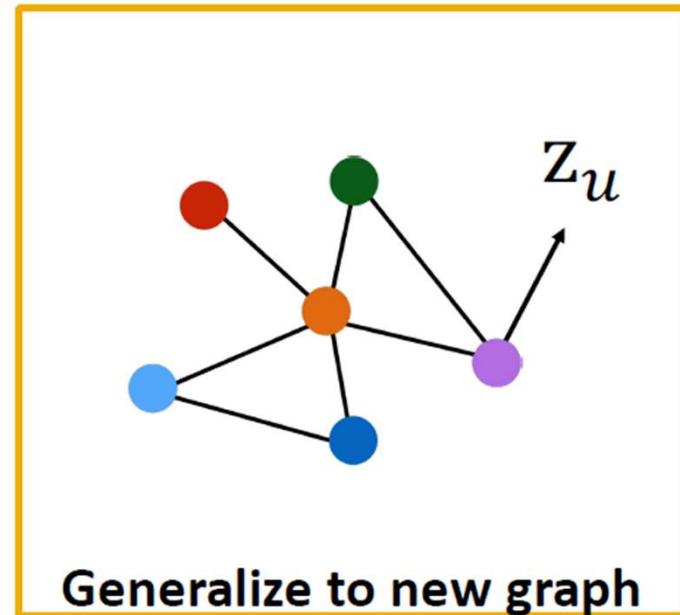
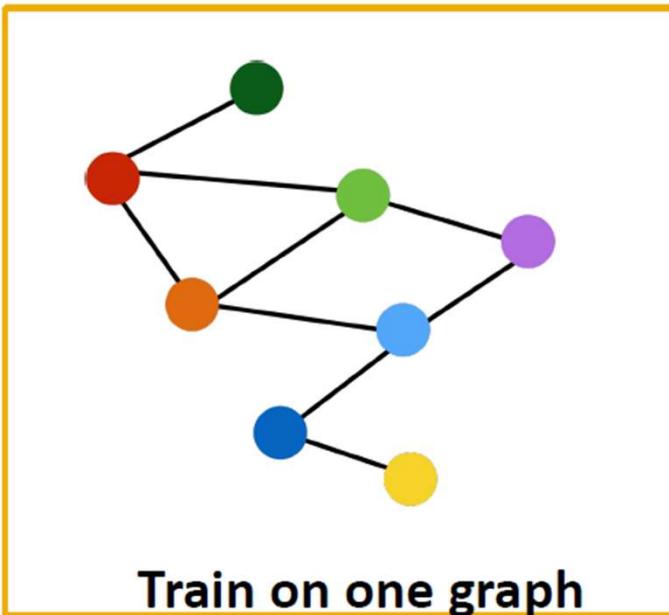


# Inductive Capability

- The same aggregation parameters are shared for all nodes:
  - The number of model parameters is sublinear in  $|V|$  and we can **generalize to unseen nodes!**



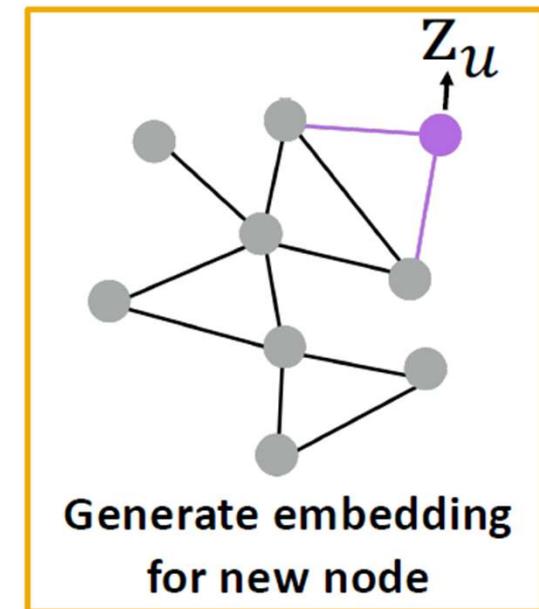
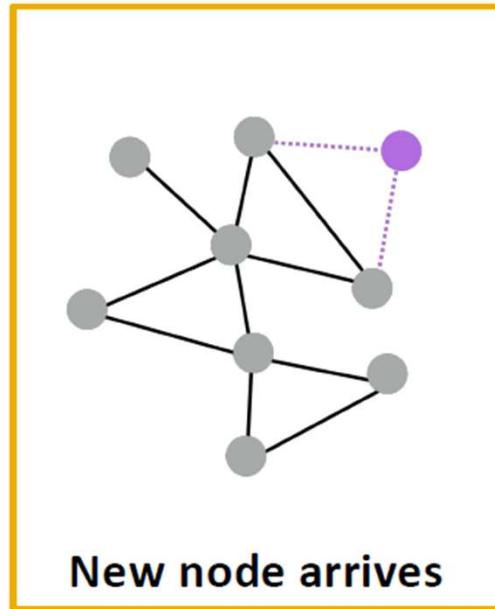
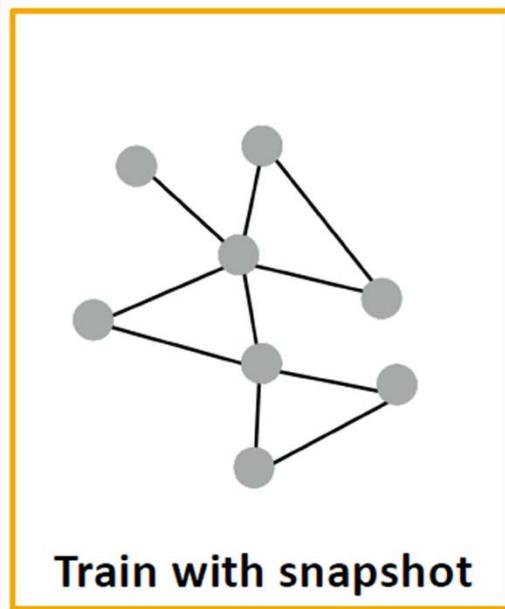
# Inductive Capability: New Graphs



Inductive node embedding → Generalize to entirely unseen graphs

E.g., train on protein interaction graph from model organism A and generate embeddings on newly collected data about organism B

# Inductive Capability: New Nodes



- Many application settings constantly encounter previously unseen nodes:
  - E.g., Reddit, YouTube, Google Scholar
- Need to generate new embeddings “on the fly”

# References

- Social Network Analysis Tanmoy Chakraborty. Chapter 9