



AIL 7022: Reinforcement Learning

Lecture 2: Hidden Markov Models

Instructor: Raunak Bhattacharyya




ScAI

YARDI SCHOOL OF ARTIFICIAL INTELLIGENCE
INDIAN INSTITUTE OF TECHNOLOGY DELHI

Why HMMs?

- Uncertainty: Zone of probabilistic reasoning
- Foundational material towards MDPs
- Constructs: Sequences of states, a.k.a. trajectories
- Algorithms: Iterative approaches
- Using observed data to make inferences

Markov Chain




Markov Chain

$\mathcal{S} = \{s_1, s_2, \dots, s_N\}$ A set of N states

$T = \begin{pmatrix} p_{11} & p_{12} & \dots & p_{1N} \\ p_{21} & p_{22} & \dots & p_{2N} \\ \vdots & \vdots & \ddots & \vdots \\ p_{N1} & p_{N2} & \dots & p_{NN} \end{pmatrix}$ A transition probability matrix

$\rho = \{p(s_1), p(s_2), \dots, p(s_N)\}$ Initial state distribution

$$p(s_i = a \mid s_1, s_2, \dots, s_{i-1}) = p(s_i = a \mid s_{i-1})$$




Markov Chain

Markov Property

Exercise

→ Compute the probability of the sequences

- ◆ Cold, Cold, Cold, Cold
- ◆ Cold, Hot, Cold, Hot



What information is missing from this question?

Initial State Distribution

Hidden Markov Model

$\mathcal{O} = \{o_1, o_2, \dots, o_M\}$ A set of M possible observations

$B = \begin{pmatrix} b_{11} & b_{12} & \dots & b_{1M} \\ b_{21} & b_{22} & \dots & b_{2M} \\ \vdots & \vdots & \ddots & \vdots \\ b_{N1} & b_{N2} & \dots & b_{NM} \end{pmatrix}$ Observation probability matrix, where $b_{ij} = p(o_j \mid s_i)$

$$\mathcal{S} = \{s_1, s_2, \dots, s_N\}$$

$$T = \begin{pmatrix} p_{11} & p_{12} & \dots & p_{1N} \\ p_{21} & p_{22} & \dots & p_{2N} \\ \vdots & \vdots & \ddots & \vdots \\ p_{N1} & p_{N2} & \dots & p_{NN} \end{pmatrix}$$

$$\rho = \{p(s_1), p(s_2), \dots, p(s_N)\}$$

$$p(o_i \mid s_1, \dots, s_i, \dots, s_T, o_1, \dots, o_i, \dots, o_T) = p(o_i \mid s_i)$$

Output Independence

Hidden Markov Model

$$O = \{o_1, o_2, \dots, o_T\}$$

**Input to the HMM: A sequence
of T observations**

$$\mathcal{S} = \{s_1, s_2, \dots, s_N\}$$

$$T = \begin{pmatrix} p_{11} & p_{12} & \dots & p_{1N} \\ p_{21} & p_{22} & \dots & p_{2N} \\ \vdots & \vdots & \ddots & \vdots \\ p_{N1} & p_{N2} & \dots & p_{NN} \end{pmatrix}$$

$$\rho = \{p(s_1), p(s_2), \dots, p(s_N)\}$$

$$\mathcal{O} = \{o_1, o_2, \dots, o_M\}$$

$$B = \begin{pmatrix} b_{11} & b_{12} & \dots & b_{1M} \\ b_{21} & b_{22} & \dots & b_{2M} \\ \vdots & \vdots & \ddots & \vdots \\ b_{N1} & b_{N2} & \dots & b_{NM} \end{pmatrix}$$

Where are HMMs used?

- Speech Recognition

Acoustic Signal

Phenomes: Pat/Bat

- Activity Recognition

Sensor Readings

Activity: walking

- Music Transcription

Audio features

Musical notes

- Finance

Financial indicators

Bull, Bear, Stable

HMM: Three Fundamental Problems

Problem 1 (Likelihood):

Given an HMM $\lambda = (\mathcal{T}, \mathcal{B})$ and an observation sequence O , determine the likelihood $P(O \mid \lambda)$.

Problem 2 (Decoding):


Given an observation sequence O and an HMM $\lambda = (\mathcal{T}, \mathcal{B})$, discover the best hidden state sequence.

Problem 3 (Learning):

Given an observation sequence O and the set of states in the HMM, learn the HMM parameters \mathcal{T} and \mathcal{B} .


Running Example: Weather and Ice Cream

- Given a sequence of observations O (each an integer representing the number of ice creams eaten on a given day) find the 'hidden' sequence of weather states (H or C)




Problem 1: Likelihood Computation

Likelihood Computation



→ Given the ice-cream eating HMM, what is the probability of the sequence of ice creams eaten being 3, 1, 3?

◆ 3 ice creams on day 1, 1 on day 2, and 3 on day 3



$p(3, 1, 3)?$


Marginalise over Hidden State Seq.

$$p(O) = p(o_1, o_2, \dots, o_T) \quad p(S) = p(s_1, s_2, \dots, s_T)$$

$$p(O) = \sum_S p(O, S)$$

$$= \sum_S p(O \mid S) \cdot p(S)$$

$$p(O \mid S) = \prod_{i=1}^T p(o_i \mid s_i)$$



Known Hidden State Seq.


$$p(3, 1, 3 \mid H, H, C)$$

$$= p(o_1 = 3, o_2 = 1, o_3 = 3 \mid s_1 = H, s_2 = H, s_3 = C)$$

$$= p(o_1 = 3 \mid s_1 = H, s_2 = H, s_3 = C) \cdot p(o_2 = 1, o_3 = 3 \mid o_1 = 3, s_1 = H, s_2 = H, s_3 = C)$$


$$= p(o_1 = 3 \mid s_1 = H) \cdot p(o_2 = 1, o_3 = 3 \mid o_1 = 3, s_1 = H, s_2 = H, s_3 = C)$$

$$p(3, 1, 3 \mid H, H, C) = p(3 \mid H) \cdot p(1 \mid H) \cdot p(3 \mid C)$$



$$p(O \mid S) = \prod_{i=1}^T p(o_i \mid s_i)$$

Marginalise via Enumeration



$$p(O) = \sum_S p(O, S)$$

$$p(3, 1, 3) = p(3, 1, 3, C, C, C) + p(3, 1, 3, C, C, H) + p(3, 1, 3, C, H, H) + \dots$$

$$= p(3, 1, 3 \mid C, C, C) \cdot p(C, C, C) + p(3, 1, 3 \mid C, C, H) \cdot p(C, C, H) + \dots$$

$$p(3, 1, 3 \mid C, C, C) = p(3 \mid C) \cdot p(1 \mid C) \cdot p(3 \mid C)$$

$$p(C, C, C) = p(s_i = C) \cdot p(C \mid C) \cdot p(C \mid C)$$

Brute Force Enumeration

Algorithm 1 Brute Force Likelihood

- 1: Enumerate all possible hidden state sequences (N^T of them)
 - 2: **for** each hidden state sequence **do**
 - 3: Compute $p(o_{1:T} \mid s_{1:T})p(s_{1:T})$
 - 4: **end for**
 - 5: Add up all the above obtained $p(o_{1:T}, s_{1:T})$ to marginalize over all possible hidden state sequences
-

Do you see a problem with this approach?

Forward Algorithm

$$\alpha_t(j) = p(o_1, o_2, \dots, o_t, s_t = j)$$

1. Initialization:

$$\alpha_1(j) = p(s_1 = j) \cdot p(o_1 \mid s_1 = j) \quad 1 \leq j \leq N$$

2. Recursion:

$$\alpha_t(j) = \sum_{i=1}^N \alpha_{t-1}(i) \cdot p(s_t = j \mid s_{t-1} = i) \cdot p(o_t \mid s_t = j) \quad 1 \leq j \leq N, 1 < t \leq T$$

3. Termination:

$$p(O) = \sum_{i=1}^N \alpha_T(i) \sum_{i=1}^N p(o_1, \dots, o_T, s_T = i)$$

**Marginalise over final
state**

Forward Algo: Rationale

$$\alpha_t(j) = p(o_1, o_2, \dots, o_t, s_t = j)$$

$$= \sum_{i=1}^N p(o_1, o_2, \dots, o_{t-1}, s_{t-1} = i, o_t, s_t = j)$$


$$= \sum_{i=1}^N p(o_{1:t-1}, s_{t-1} = i) \cdot p(o_t, s_t = j \mid o_{1:t-1}, s_{t-1} = i)$$

$$= \sum_{i=1}^N p(o_{1:t-1}, s_{t-1} = i) \cdot p(s_t = j \mid o_{1:t-1}, s_{t-1} = i) \cdot p(o_t \mid o_{1:t-1}, s_{t-1} = i, s_t = j)$$

$$= \sum_{i=1}^N p(o_{1:t-1}, s_{t-1} = i) \cdot p(o_t \mid s_t = j) \cdot p(s_t = j \mid s_{t-1} = i)$$

$$\alpha_t(j) = \sum_{i=1}^N \alpha_{t-1}(i) \cdot p(s_t = j \mid s_{t-1} = i) \cdot p(o_t \mid s_t = j)$$

Brute Force vs. Forward Algorithm



$$\begin{aligned} p(o_{1:T}) &= \sum_{\forall \{s_{1:T}\}} p(o_{1:T}, s_{1:T}) \\ &= \sum_{\forall \{s_{1:T}\}} p(o_{1:T} | s_{1:T}) \cdot p(s_{1:T}) \\ &= \sum_{\forall \{s_{1:T}\}} \left(\prod_{i=1}^T p(o_i | s_i) \cdot p(s_{1:T}) \right) \end{aligned}$$

$$\begin{aligned} p(o_{1:T}) &= \sum_{i=1}^N p(o_{1:T}, s_T = s^i) \\ &= \sum_{i=1}^N p(o_{1:T}, s_T = i) \end{aligned}$$

Notation alert

Forward Algorithm

$$\alpha_t(j) = p(o_1, o_2, \dots, o_t, s_t = j)$$

1. Initialization:

$$\alpha_1(j) = p(s_1 = j) \cdot p(o_1 \mid s_1 = j) \quad 1 \leq j \leq N$$

2. Recursion:

$$\alpha_t(j) = \sum_{i=1}^N \alpha_{t-1}(i) \cdot p(s_t = j \mid s_{t-1} = i) \cdot p(o_t \mid s_t = j) \quad 1 \leq j \leq N, 1 < t \leq T$$

3. Termination:

$$p(O) = \sum_{i=1}^N \alpha_T(i)$$

How do we translate this to pseudocode?

Forward Algo: Pseudocode

$$\alpha_t(j) = p(o_1, o_2, \dots, o_t, s_t = j)$$

function FORWARD(*observations* of len T , *state-graph* of len N) **returns** *forward-prob*

create a probability matrix *forward*[N, T]

for each state s **from** 1 **to** N **do**

$forward[s, 1] \leftarrow \pi_s * b_s(o_1)$

for each time step t **from** 2 **to** T **do**

for each state s **from** 1 **to** N **do**

$$forward[s, t] \leftarrow \sum_{s'=1}^N forward[s', t-1] * a_{s', s} * b_s(o_t)$$

$$forwardprob \leftarrow \sum_{s=1}^N forward[s, T]$$



return *forwardprob*

$$\alpha_1(j) = p(s_1 = j) \cdot p(o_1 \mid s_1 = j)$$

$$\alpha_t(j) = \sum_{i=1}^N \alpha_{t-1}(i) \cdot p(s_t = j \mid s_{t-1} = i) \cdot p(o_t \mid s_t = j)$$

$$p(O) = \sum_{i=1}^N \alpha_T(i)$$

Forward Trellis



How do you compute the original goal: observation sequence likelihood?

HMM: Three Fundamental Problems

Problem 1 (Likelihood):



Given an HMM $\lambda = (\mathcal{T}, \mathcal{B})$ and an observation sequence O , determine the likelihood $P(O \mid \lambda)$.

Problem 2 (Decoding):


Given an observation sequence O and an HMM $\lambda = (\mathcal{T}, \mathcal{B})$, discover the best hidden state sequence.

Problem 3 (Learning):


Given an observation sequence O and the set of states in the HMM, learn the HMM parameters \mathcal{T} and \mathcal{B} .

Problem 2: Decoding

Problem 2: Decoding



→ Find the hidden state sequence that was most likely to have generated the input observation sequence.



$p(3, 1, 3)?$

Earlier goal

What should the goal be now?

$$\arg \max_{s_1, s_2, s_3} p(s_1, s_2, s_3 \mid o_1 = 3, o_2 = 1, o_3 = 3)?$$

Decoding: Equivalent Objective

$$\arg \max_{s_1, s_2, s_3} p(s_1, s_2, s_3 \mid o_1 = 3, o_2 = 1, o_3 = 3)?$$

$$p(s_1, s_2, s_3 \mid o_1, o_2, o_3) = \frac{p(o_1, o_2, o_3, s_1, s_2, s_3)}{p(o_1, o_2, o_3)}$$

$$\arg \max_{s_1, s_2, s_3} p(s_1, s_2, s_3 \mid o_1, o_2, o_3) = \arg \max_{s_1, s_2, s_3} p(o_1, o_2, o_3, s_1, s_2, s_3)$$

Where have we seen this quantity before?

$$p(o_{1:3}, s_{1:3})$$

$$p(o_{1:3}) = \sum_{\forall \{s_{1:3}\}} p(o_{1:3}, s_{1:3})$$

Can you think of a brute force algorithm for decoding?

Viterbi Algorithm: Building Blocks

$$\arg \max_{s_1, s_2, s_3} p(o_1, o_2, o_3, s_1, s_2, s_3)$$

Define: $v_t(j) = \max_{s_1, s_2, \dots, s_{t-1}} p(s_1, s_2, \dots, s_{t-1}, o_1, o_2, \dots, o_t, s_t = j)$

$$v_t(j) = \max_{s_{1:t-1}} p(s_{1:t-1}, o_{1:t}, s_t = j)$$

How do we use this to get what we want?

Goal: $\max_{s_1, s_2, s_3} p(s_1, s_2, s_3, o_1, o_2, o_3)$

$$\max_{s_1, s_2} p(s_1, s_2, o_1, o_2, o_3, s_3 = C)$$

$$\max_{s_1, s_2} p(s_1, s_2, o_1, o_2, o_3, s_3 = H)$$

$$\max_{i=1}^N \max_{s_1, s_2} p(s_1, s_2, o_1, o_2, o_3, s_3 = i)$$

Goal obtained by: $\max_{i=1}^N v_T(i)$

Viterbi Algo: Recursion

$$v_t(j) = \max_{s_{1:t-1}} p(s_{1:t-1}, o_{1:t}, s_t = j)$$

$$p(s_{1:t-1}, o_{1:t}, s_t = j) = p(s_{1:t-2}, o_{1:t-1}, s_{t-1}, o_t, s_t = j)$$

$$= p(s_{1:t-2}, o_{1:t-1}, s_{t-1}) \cdot p(o_t, s_t = j \mid s_{1:t-2}, o_{1:t-1}, s_{t-1})$$

$$= p(s_{1:t-2}, o_{1:t-1}, s_{t-1}) \cdot p(s_t = j \mid s_{1:t-2}, o_{1:t-1}, s_{t-1}) \cdot p(o_t \mid s_t = j, s_{1:t-2}, o_{1:t-1}, s_{t-1})$$

$$= p(s_{1:t-2}, o_{1:t-1}, s_{t-1}) \cdot p(s_t = j \mid s_{t-1}) \cdot p(o_t \mid s_t = j)$$

$$p(s_{1:t-1}, o_{1:t}, s_t = j) = p(s_{1:t-2}, o_{1:t-1}, s_{t-1}) \cdot p(s_t = j \mid s_{t-1}) \cdot p(o_t \mid s_t = j)$$

Viterbi Algo: Recursion

$$p(s_{1:t-1}, o_{1:t}, s_t = j) = p(s_{1:t-2}, o_{1:t-1}, s_{t-1}) \cdot p(s_t = j \mid s_{t-1}) \cdot p(o_t \mid s_t = j)$$

$$\max_{s_{1:t-1}} p(s_{1:t-1}, o_{1:t}, s_t = j)$$

$$= \max_{i=1}^N \max_{s_{1:t-2}} p(s_{1:t-2}, o_{1:t-1}, s_{t-1} = i) \cdot p(s_t = j \mid s_{t-1} = i) \cdot p(o_t \mid s_t = j)$$

$$v_t(j) = \max_{i=1}^N v_{t-1}(i) \cdot t_{ij} \cdot b_j(o_t)$$

Viterbi Algorithm

1. Initialization:

$$\begin{aligned}v_1(j) &= \pi_j b_j(o_1) & 1 \leq j \leq N \\bt_1(j) &= 0 & 1 \leq j \leq N\end{aligned}$$

2. Recursion

$$\begin{aligned}v_t(j) &= \max_{i=1}^N v_{t-1}(i) a_{ij} b_j(o_t); & 1 \leq j \leq N, 1 < t \leq T \\bt_t(j) &= \operatorname{argmax}_{i=1}^N v_{t-1}(i) a_{ij} b_j(o_t); & 1 \leq j \leq N, 1 < t \leq T\end{aligned}$$

3. Termination:

$$\text{The best score: } P^* = \max_{i=1}^N v_T(i)$$

$$\text{The start of backtrace: } q_T^* = \operatorname{argmax}_{i=1}^N v_T(i)$$

Viterbi Algorithm: Pseudocode

create a path probability matrix $viterbi[N,T]$

for each state s from 1 to N do

$viterbi[s,1] \leftarrow \pi_s * b_s(o_1)$

$backpointer[s,1] \leftarrow 0$

$$v_1(j) = \pi_j b_j(o_1) \quad 1 \leq j \leq N$$

$$bt_1(j) = 0 \quad 1 \leq j \leq N$$

for each time step t from 2 to T do

; recursion step

for each state s from 1 to N do

$$viterbi[s,t] \leftarrow \max_{s'=1}^N viterbi[s',t-1] * a_{s',s} * b_s(o_t)$$

$$backpointer[s,t] \leftarrow \operatorname{argmax}_{s'=1}^N viterbi[s',t-1] * a_{s',s} * b_s(o_t)$$

$$v_t(j) = \max_{i=1}^N v_{t-1}(i) a_{ij} b_j(o_t);$$

$$bt_t(j) = \operatorname{argmax}_{i=1}^N v_{t-1}(i) a_{ij} b_j(o_t)$$

$$bestpathprob \leftarrow \max_{s=1}^N viterbi[s,T]$$

; termination step

$$bestpathpointer \leftarrow \operatorname{argmax}_{s=1}^N viterbi[s,T]$$

; termination step

$bestpath \leftarrow$ the path starting at state $bestpathpointer$, that follows $backpointer[]$ to states back in time

$$P^* = \max_{i=1}^N v_T(i)$$

$$q_T^* = \operatorname{argmax}_{i=1}^N v_T(i)$$

Viterbi Trellis

